

1 **Understanding AI Data Production & Community Impacts Worldwide: A
2 Multivocal Literature Review**

3 ANONYMOUS AUTHOR(S)

4 Artificial intelligence (AI) depends on data production: the sociotechnical process that transforms human knowledge into computational
5 resources. The connections among AI systems, data practices, and impacts on Indigenous, underrepresented, and underserved
6 communities—though critical—have not been systematically examined. To this end, we conduct a Multivocal Literature Review (MLR)
7 integrating 350 academic and grey-literature sources to analyze how AI systems, data practices, and community impacts intersect.
8 Across five analytic domains—Data Relations, Data Labor, Data Representation, Data Infrastructure, and Data Governance—we
9 distinguish extractive data production mechanisms that centralize control from high-agency pathways in which communities exercise
10 authority. We contribute (1) a multivocal review that positions data production as a site of sociotechnical power rather than a technical
11 prerequisite; (2) implications for HCI research including upstream infrastructure as a design site, provenance-first architectures, and
12 federated data governance supporting community sovereignty; (3) methodological validation of multivocal synthesis for bridging
13 academic critique with community practice; and (4) an open corpus mapping sources across pipeline stages, historical eras, and
14 geographic contexts.

15 CCS Concepts: • **Do Not Use This Code → Generate the Correct Terms for Your Paper**; *Generate the Correct Terms for Your
16 Paper*; Generate the Correct Terms for Your Paper; Generate the Correct Terms for Your Paper.

17 Additional Key Words and Phrases: AI, ML pipeline, data production, extractive practices, underserved communities, Indigenous data
18 sovereignty, data collection

19 **ACM Reference Format:**

20 Anonymous Author(s). 2018. Understanding AI Data Production & Community Impacts Worldwide: A Multivocal Literature Review.
21 In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (Conference acronym 'XX)*. ACM,
22 New York, NY, USA, 33 pages. <https://doi.org/XXXXXXX.XXXXXXX>

23 **1 Introduction**

24 Contemporary artificial intelligence (AI) systems depend on data. As approaches have advanced over the past three
25 decades, the scale and composition of data needs has transformed: from small expert-curated datasets like MNIST [83],
26 to massive crowdsourced benchmarks such as ImageNet (Deng et al., 2009), and now to foundation models trained
27 on billions of scraped web documents, images, and interaction traces [128, 142]. The opacity and complexity of the
28 ML pipeline [22], as well as the diversity and amount of human knowledge and labor needed [126, 163], has expanded
29 dramatically.

30 This trajectory matters because the choices made in gathering and curating data directly shape which communities
31 benefit from AI systems and which communities bear their costs. Data is made, not found—it is produced through
32 a series of choices about what to gather, how to curate it, and under what terms. Decision made upstream and all
33 along the pipeline can either create or mitigate harms, which fall disproportionately on underserved, underrepresented,

34 Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not
35 made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components
36 of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on
37 servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

38 © 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.

39 Manuscript submitted to ACM

40 Manuscript submitted to ACM

and Indigenous communities [7, 13, 78, 87, 100, 123]. Yet despite growing critical attention to algorithmic harms and dataset bias, data production itself—the complex sociotechnical process through which human knowledge becomes computational input—has rarely been treated as a central object of inquiry within HCI.

Relevant literature is scattered across disciplines and publication ecosystems. Human-computer interaction (HCI) contributes a long-standing body of research on how sociotechnical systems enact power, from canonical postcolonial critiques [66, 154] to recent work analyzing “extractive” dynamics in ICT4D research [45], and epistemic injustice [5, 168]. Review literature in HCI and adjacent fields offers substantial insight into how AI systems affect communities. A scoping review by Shelby et al. [146] maps harms experienced by underserved groups, and a subsequent systematic review by Wang et al. [165] builds on this foundation through a focused synthesis across disability contexts. Case-based analyses examine how AI systems mismatch their contexts of use ([136]. Surveys of AI ethics and NLP bias identify structural gaps and the absence of lived-experience perspectives [16, 18], while dataset genealogies trace how collection and annotation practices embed exclusions [39]. Recent work analyzes how misabstraction cascades through sociotechnical systems [35] and takes stock of social-justice commitments within HCI [27].

This body of work illuminates important dimensions of a three-part relationship: AI systems, data production practices, and community impacts. Yet, most scholarship examines discrete pieces rather than synthesizing across all three. Most reviews also draw primarily on academic publications, capturing the scientific state-of-the art but leaving less visible the state-of-practice documented in policy reports, organizational materials, and community outputs. These gaps call for synthesis across this three-part relationship. To this end, we use a multivocal literature review (MLR), an approach designed to synthesize knowledge that circulates across many publication ecosystems. An MLR integrates peer-reviewed research with policy documents, organizational reports, technical specifications, and community-generated materials (“white” and “grey” literature), and as such offer a structured way to work with a wider evidence landscape [51]; see Appendix A.

We treat *data production* as the complex sociotechnical process through which data is defined, gathered, curated, and controlled across model pipelines.¹ This framing supports a move beyond understanding *data collection* as a routine methodological disclosure or neutral technical artifact. Instead, we center the institutional choices, power relations, and consequences that underpin AI development and determine who benefits from or bears its costs. Our approach leverages Critical Computing as a diagnostic lens and Social Justice as a normative orientation. Critical Computing shows how data practices reflect institutional priorities and labor arrangements rather than some objective “ground truth” modeled by engineers, and offers tools for analyzing power in dataset construction and use. Social Justice complements the diagnosis by asking how data work might redistribute agency, benefit, and governance toward the communities whose knowledge and labor support AI systems. Together, the two orientations clarify why upstream data production is a sociotechnical domain of timely concern for HCI and a site for upstream design intervention.

Overall, this work provides a literature review with novel concepts of data production as a sociotechnical site where power is negotiated and encoded, rather than a logistical preliminary to model development. We identify mechanisms where extractive practices and high-agency alternatives diverge across the ML pipeline, with implications for HCI. In summary, we contribute:

- A multivocal literature review of 350 sources across academic and grey literature, resulting in five analytic domains where AI systems, data production, and community impacts intersect

¹We share with Miceli & Posada [96] an emphasis on “production” to foreground relations of power and knowledge in data and labor, which echoes the “assemblage” approach of Kitchin et al. [76], also rooted in Foucauldian critique.

- 105 • Opportunities for HCI research and practice, including upstream data infrastructure as a design site, provenance-
106 first architectures, federated learning for community sovereignty, and ethics review paradigms that scrutinize
107 data production
- 109 • Methodological validation of multivocal synthesis for sociotechnical inquiry, showing how grey literature
110 captures state-of-practice missing from academic venues;
- 112 • An open corpus of 350 sampled sources mapped across pipeline stages, historical eras, and geographic contexts,
113 with structured summaries and documented rationales for relevance to the three-part inquiry, publicly available
114 at <https://github.com/ADC-chi/ai-data-production-landscape>.

116 1.1 Key Terms and Definitions

118 1.1.1 *Artificial intelligence.* For clarity, we use “AI” in this paper primarily in its modern ML sense. An ML system learns
119 patterns and rules from training data to create a predictive model [22]. The resulting model must then be evaluated
120 for reliability and generalization using a separate, independent dataset known as test data [42]. We contextualize our
121 discussion of AI within its broader historical trajectory of research and development but focus on the current statistical
122 and data-driven era of AI that facilitates many contemporary extractive regimes [55]. We intend our discussion to be
123 situated not merely as a critique of modern ML but as a reflection on a continuous thread within technological and
124 social history.

127 1.1.2 *Extractive.* We use “extractive” in this paper to denote high-asymmetry or dispossessive practices, building upon
128 its conventional association with Indigenous marginalization and digital forms of resource appropriation. This definition
129 is designed to encompass diverse historical and contemporary manifestations of power imbalances that result in one
130 party’s advantage at the expense of another’s autonomy or resources [19, 114]. Illustrative examples of such high-
131 asymmetry practices are evident in historical contexts, such as the exploitative labor practices of UK coal mining [33];
132 the profoundly unethical nature of the Tuskegee syphilis experiment in the United States and the untreated carcinoma
133 study in New Zealand [70, 121]; and contemporary issues like large-scale industrial mining [107] and pervasive AI
134 surveillance [118]. By broadening this definition, our objective is to more accurately encapsulate the systemic character
135 of extraction across various domains. In contrast, we describe as high-agency examples of principles and practices in
136 the literature which prioritize active participation, equitable distribution of power, and community-defined obligations
137 [124?]. These examples appear in contexts where communities, practitioners, or institutions negotiate shared authority,
138 shape the terms of data contribution, or establish governance arrangements that align data use with locally grounded
139 priorities.

144 1.1.3 *Communities and populations.* We use “underserved” to describe communities lacking adequate infrastructural,
145 institutional, or economic support, and underrepresented to indicate groups whose knowledge, languages, or perspectives
146 are numerically absent or devalued in AI research and development [88, 146]. We use the umbrella category “Indigenous,”
147 which “enables historically and geographically separated peoples to recognize each other and their common plight, and
148 to collaborate towards a better future” [133]. We avoid “marginalized” in the adjectival form to emphasize agency and
149 resistance rather than positioning communities as passive victims. We use Global Majority to emphasize that most of
150 the world’s population lies outside Euro-American contexts. Our chosen terms underscore structural asymmetries in
151 power and resource distribution rather than deficits within communities themselves [155].

157 **2 Background**

158 **2.1 Critical Traditions on Extraction and Justice**

160 Foundational works from theoretical, historical, and community traditions establish frameworks for studying power in
161 knowledge production. Theories of epistemic violence and injustice [47, 150] and “situated knowledges” [62] interrogate
162 how knowledge systems encode relations of domination. Historical analyses of colonial resistance show alternative
163 epistemologies and organizing strategies [68].

165 Black feminist theory articulates intersectional approaches to structural power, from early collective statements [30]
166 to analyses of interlocking systems of oppression [31]. Gender and queer theory establish frameworks for analyzing
167 the production of normativity [23], binary logics [144], and classificatory power [29]. Indigenous studies center
168 community sovereignty and relational ethics [9, 37, 86] and provide frameworks for decolonizing knowledge production
169 in research [149] and AI data practices [20]. Critical data studies crystallize a complementary set of concerns for the
170 digital context, with a focus on datafication, surveillance, and governance [77].

172 Foundational works attune us to centuries of extractive patterns, resistance, and knowledge-making. They are
173 essential for understanding present and future technological worlds. Here, critical works anchor the conceptual
174 vocabulary of extraction and justice in the context of the global AI data production ecosystem.

178 **2.2 Evolution of AI Data Production Practices**

180 Since the advent of machine learning, there has been a constant need for data. Over time, how that data was produced
181 has undergone transformations beyond dataset sizes. These changes include how data is produced and who performs
182 the work [39]. As demands for larger models have intensified, practices have shifted from small, carefully curated
183 corpora, to large datasets assembled through web-scraping and crowdsourced annotations, to massive, automated
184 web-scraped collections supported by industrial-scale annotation.

186 *Era 1.* Early curated datasets were small, domain-specific, and selected by experts. A canonical example is MNIST, a
187 dataset of handwritten digits drawn from U.S. postal codes [83]. Choices about inclusion and categorization reflected
188 institutional knowledge and disciplinary priorities.

189 *Era 2.* Large curated datasets expanded scale through crowdsourced annotation of web-scraped content, exemplified
190 by ImageNet [38] and MS COCO [90]. This era accelerated deep learning [36, 56] but shifted labor from domain
191 experts to distributed workers, often in Global Majority regions, as well as automated web-scraping efforts, ultimately
192 emphasizing performance gains over contextual fit.

193 *Era 3.* Contemporary data production diverges into two parallel approaches. Massive, largely uncurated web-scraped
194 corpora such as Common Crawl, C4 [41, 130], LAION [142, 143], Refined Web [122], and ClueWeb22 [113]) are assembled
195 through automated scraping at an unprecedented scale. Such production efforts shape and are shaped by competitive
196 foundation model development [15, 153]. Alongside and often in response, smaller, highly curated datasets emerged,
197 produced through participatory methods and community partnerships. Examples include ROOTS [82], Masakhane’s
198 African language collections [109], and Cohere’s multilingual Aya Dataset [148].

199 Dataset hosting and governance practices have shifted over time as well: from freely downloadable units like MNIST,
200 to single-location storage on cloud services (e.g., AWS, HuggingFace), to URL-indexed collections like LAION that
201 disclaim responsibility for original sources, and to emerging federated “data spaces” designed to support locally owned
202 infrastructures and community governance [64]. Proprietary datasets are closed, and open-source alternatives range

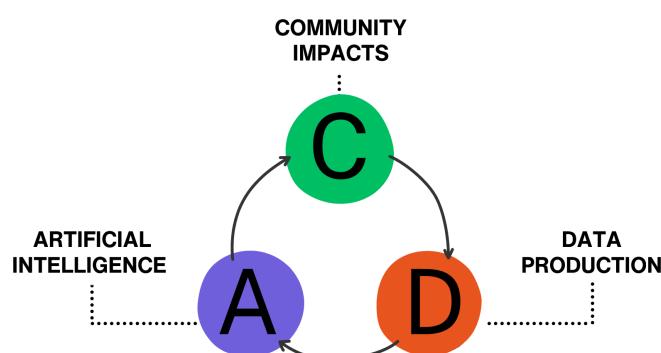


Fig. 1. A/D/C scoping framework for identifying relevant literature across three interrelated domains of AI (A), Data Production (D), and Community Impacts (C).

from massive scrapes to carefully stewarded community collections; each modality comes with distinct risks and obligations [89, 169].

The three eras feature distinct technical capabilities, institutional arrangements, and methodologies that enabled extractive patterns to scale industrially. While both eras 2 and 3 are founded in web-scraped datasets, the unprecedented scale at which era 3 extracts data is unprecedented and as such, the current era hosts both the most expansive extractive practices and the most developed community-controlled frameworks. The future of AI data production is not determined.

3 Methodology

We conducted an MLR to assemble a structured body of evidence on AI data production and its impacts on underserved, underrepresented, and Indigenous communities. We treated academic and grey sources as complementary evidence streams. Figure 1 presents the research design and methodology.

Because our inquiry has three interrelated parts—AI systems, data production practices, and community impacts—we formalized it through what we call the A/D/C framework (Figure 1). This framework defined the scope of our inquiry, which was ultimately bounded by community impacts (C). All sources engage substantively with community experiences, power dynamics, or consequences in contexts relevant to AI data production. Sources varied in whether they directly addressed AI systems (A) and data production practices (D), or provided foundational understanding that informs interpretation of these dimensions. Section 3.2 provides more detail on how we created the corpus based on this framework.

3.1 Search Strategy

We developed the search strategy by deriving keywords from the A/D/C framework. We searched seven academic databases between August 2024 and January 31 2025: ACM Digital Library, IEEE Xplore, ScienceDirect, Taylor & Francis Online, Wiley Online Library, Springer Link, and Google Scholar. We iteratively developed boolean search strings with AND/OR terms across variants of A/D/C terms using Boolean operators. Titles and abstracts were screened first, followed by full-text assessment for items meeting initial criteria. Table 1 shows the key and supplementary search terms of our inquiry.

Table 1. Key and Supplementary Search Terms

Dimension	Key Terms	Supplementary Terms
AI Systems (A)	Artificial Intelligence, Machine Learning, AI, ML	Large Language Models, LLMs, Computer Vision, Foundation Models, Neural Networks, Automated Systems, Algorithmic Systems, Deep Learning
Data Production (D)	Data Collection, Data Production, Dataset Creation, Data Curation, Data Practices	Annotation, Labeling, Data Labor, Crowdsourcing, Web Scraping, Data Extraction, Dataset Development, Data Work, Data Gathering, Responsible AI
Community Impacts (C)	Indigenous, Marginalized, Underrepresented, Underserved, Community	Global Majority, Global South, Data Sovereignty, Linguistic Diversity, Cultural Context, Extraction, Appropriation, Bias, Fairness, Harm, Safety

We developed four primary query sets: foundational (targeting core data production practices in AI contexts affecting communities), extraction frame (targeting exploitative practices), data labor (focusing on crowdsourcing and platform labor), and alternatives (seeking participatory and community-led approaches). The ACM Digital Library search illustrates our results. Across the four query sets, 1,914 hits yielded 1,201 items screened, 153 meeting initial criteria, and 48 unique sources after full-text review and duplicate removal. Similar strategies applied to the remaining six databases. Database searches contributed 174 sources, representing 50% of the final corpus. See B for more search details.

For grey literature we used different methods. Following Garousi et al. [51], we used general Google Search and systematically examined organizational ecosystems engaged in AI data work, prioritizing organizational reports, policy documents, and community outputs from established entities. Three complementary methods supplemented database and grey literature searches. Citation snowballing [170] from 20 seed papers tracked forward and backward citations iteratively, contributing 51 sources (15%). Hand-searching of journals including CHI, FAccT, CSCW, *Journal on Responsible Computing*, ACL, *Big Data & Society*, and *AI & Society* contributed 21 sources (6%). Iterative gap-filling searches addressed underrepresented regions, concepts, or pipeline stages as the corpus took shape, contributing 31 sources (9%).

Inclusion criteria followed from the A/D/C framework. Sources entered the corpus when they engaged community impacts substantively. Most sources additionally provided direct evidence about data production practices or AI systems. This meant including sources that analyzed AI system behavior, deployment, or evaluation in relation to community outcomes; examined data sourcing, processing, annotation, governance, or infrastructure with implications for affected communities; provided community-governed protocols, sovereignty statements, or governance frameworks; or established theoretical or epistemological foundations addressing power, resistance, extraction, or marginalization in ways essential for interpreting AI data practices and their consequences. We excluded sources that, for example, discussed AI ethics, fairness, or responsible AI at a high level without addressing data practices or community impacts; focused solely on model performance, technical optimization, or algorithmic advances without sociotechnical analysis; reported community-based research unrelated to AI systems or data production; or were non-English (a pragmatic limitation we discuss more below).

Quality assessment varied by source type and followed guidance from Kamei et al. [72]. Academic items underwent venue peer review. Grey-literature items required additional evaluation; we assessed organizational authority, author expertise, community recognition, and the provenance of policy documents, and we interpreted community outputs through alignment with decolonizing methodologies and community endorsement. These criteria follow guidance for

grey-literature appraisal in multivocal reviews and draw on elements of Garousi et al. [51]’s framework, including stated aims, methodological clarity, contribution, and outlet type—adapted for community-governed and sovereignty-oriented materials.

Screening proceeded in two stages: titles and abstracts were reviewed for relevance to the C boundary, followed by full-text assessment. The first author led database and grey literature searches. Two authors independently read full-text articles, prepared summaries, and presented sources in batches to the full team for consensus review. Disagreements on inclusion criteria and relevance to the A/D/C framework were resolved through discussion. This process occurred in two rounds (August–September 2024), with each round reviewing approximately 100 candidate sources. These consensus rounds established shared standards before the two authors completed full screening of the 350-source corpus in February 2025. The final corpus contains 258 academic items (74%) and 92 grey-literature items (26%).

3.2 Corpus Creation

We created a datasheet that categorizes each source across multiple dimensions to provide maximum contextualization [35]. Coded categories included bibliographic metadata (author, year, venue, type), A/D/C coverage, pipeline stage, historical era, orientation, geographic focus, and author affiliation. For each source, we additionally recorded a unique rationale for A/D/C coverage and a brief summary to support traceability of sourcing and selection decisions. The same two authors who led source selection conducted categorization using the collaborative consensus approach established during screening.

A/D/C Coverage. We categorized each source based on where direct evidence appears across our three-part inquiry: AI systems (A), data production practices (D), and community impacts (C). Every source engages all three dimensions analytically, but sources vary in what they directly support versus what requires interpretive connection. We wrote rationales stating what each source contributes to A, D, and C to clarify where direct evidence appears and where relevance is interpretive and to provide transparent disclosure of our interpretive stance on each source. Tags indicate where direct evidence is present. We do not tag A or D dimensions alone because all sources must engage community impacts (C) to enter the corpus. Our coding produced four categories, described in Table 2.

Pipeline Stages. We coded each source to stages of a simplified AI development pipeline (Figure 2): *Problem Understanding & Formulation* (institutional prioritization, funding decisions, and product conception), *ML System Design and Development* (data selection and enrichment, model architecture choices, and training processes), and *Deployment & Impact* (product testing, launch, and post-deployment effects) (Martin, 2020). The interconnected nature of these processes means that problems at one stage “cascade” through the pipeline, compounding downstream issues and harms [134, 152]. Sources addressing multiple stages or cross-cutting dynamics were tagged accordingly.

Historical Eras. We distinguished three eras of data production, per discussion in 2.2: Era 1 (expert-curated datasets, pre-2009), Era 2 (crowdsourced benchmarks, 2009–2017), and Era 3 (web-scraped and foundation models, 2017–present). Multi-era sources were coded accordingly. No sources were coded exclusively as Era 1.

Orientations. Each source received a single orientation code reached through team consensus based on the primary analytical purpose the source served in this inquiry. *Extractive* sources provided direct evidence of practices undermining consent, compensation, or community benefit. *High-agency principles* advanced normative frameworks with explicit policy or governance recommendations. *High-agency practices* described operationalized initiatives with concrete implementation details.

Synthesis. We synthesized findings through iterative analysis across these dimensions. When multiple sources described similar mechanisms across different contexts, we consolidated these into recurring patterns. Individual

365 Table 2. Triangle coverage coding definitions showing how corpus sources engage AI systems (A), data production practices (D), and
 366 community impacts (C)

368 Code	369 Description	370 Example Sources
371 ADC	372 Direct evidence relevant 373 to AI systems, data 374 production, and 375 community impacts	376 Garcia et al. [49] on critical refusal as an intervention into extractive data 377 logics and governance; Hall et al. [57] on participatory, community-engaged dataset production; Park et al. [117] on designing accessible infrastructures for collecting AI data from people with disabilities; Rifat et al. [131] on categorization politics and context erasure in annotating faith-based violence data; Lewis et al. [86] on Indigenous protocol-aligned dataset construction and culturally grounded AI applications
378 DC	379 Direct evidence for data 380 production and 381 community impacts; AI 382 relevance is interpretive	383 Adley et al. [4] on ethical data collection with marginalized groups and 384 power dynamics in practice; Cooper et al. [32] on community-collaborative research models emphasizing shared control and benefit; Hancock et al. [58] on tensions in data sharing and harms within a modern slavery data ecosystem; Taylor and Kukutai [156] on Indigenous Data Sovereignty and metadata governance; Pool [125] on colonial census practices replacing Māori knowledge systems
385 AD	386 Direct evidence for AI 387 systems and data 388 production; community 389 impacts are clearly 390 implied	391 Bhardwaj et al. [11] on evaluating ML datasets through a data-curation 392 lens and FAIR principles; Koch et al. [80] on dataset reuse and 393 benchmark concentration; Sambasivan et al. [134] on data cascades and 394 hidden labor in high-stakes ML pipelines; Schiff et al. [141] on 395 translating AI principles into practice via participatory, iterative impact 396 assessment; Zhao et al. [173] on fairness-curation challenges faced by 397 dataset curators across organizational and socio-political contexts
398 C	399 Direct evidence about 400 community impacts only; 401 A and D relevance is 402 interpretive	403 Battiste [9] on Indigenous epistemologies and marginalization; Haraway 404 [62] on situated knowledge and partial perspective; Igwe et al. [65] on 405 non-extractive research principles; James [68] on colonial extraction 406 economies and collective resistance in the Haitian Revolution; Shapiro 407 and McNeish [145] on hyper-extractivism and resistance

400 sources could exemplify multiple patterns. Patterns were distilled into the five analytic domains described in Section 4.
 401 Complete coding definitions with examples appear in Appendix C.

404 3.3 Limitations and Reflexivity

405 We recognize that subjectivity shapes our interpretations, though transparent documentation and multi-researcher
 406 validation helped address this inherent limitation. Connecting multiple disciplinary traditions, historical eras, and global
 407 contexts proved challenging, creating “translation needs” across distinct vocabularies and epistemological frameworks.
 408 The sourcing strategy privileged networks in Africa and global Indigenous movements, yielding detailed coverage of
 409 those ecosystems. Parallel developments in Middle Eastern, Southeast Asian, and Latin American contexts appear less
 410 frequently, not because such initiatives were absent, but because they circulated in networks less accessible to our
 411 inquiry. We acknowledge that English-language search restrictions inherently reinforce Western-centric representation.
 412 Scholarly and community infrastructures condition what becomes visible in review corpora, creating unevenness
 413

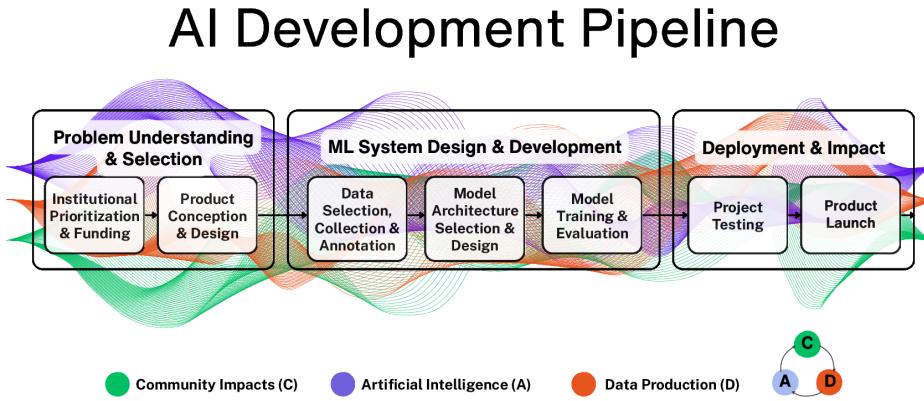


Fig. 2. Simplified AI development pipeline with three interwoven sine waves representing AI (purple), data production (orange), and community impacts (green) flowing continuously through all pipeline stages. Pipeline diagram derived from Martin [95].

despite our efforts. Therefore, we believe it is important to consider our own identities alongside the analysis of this work given that our backgrounds and perspectives may bias the interpretation of this work [34].

The authors hold diverse racial and ethnic identities (Black, White, Mixed-race) with cultural roots in the United States, Canada, South Africa, Ghana, Japan, and France. These backgrounds shaped our ability to recognize and access specific community-led networks (particularly in African and Indigenous contexts) while leaving others less visible to us. In terms of epistemic lens, we work across industry and academia, with backgrounds in computer science, HCI, and the humanities. This dual positioning allowed us to bridge the gap between technical documentation and critical theory, for example, recognizing “grey literature” as rigorous evidence of high-agency alternatives. However, our location within these professionalized research institutions also means we likely missed grassroots resistance tactics that do not circulate in written or digital forms. We recognize that we are observing extraction from within the institutions that often facilitate it, and we present these findings as a necessary, though partial, mapping of the landscape.

4 Findings

We structure our findings according to five domains of data production, which we have conceptualized as analytic elements. Rather than logistical steps, these domains function as sites where power is negotiated, contested, and encoded: **Data Relations** (the negotiation of agency and terms of engagement), **Data Labor** (the creation versus capture of value), **Data Representation** (the exercise of epistemic authority through categorization), **Data Infrastructure** (the allocation of capacity and provenance), and **Data Governance** (the enforcement of sovereignty and accountability). Within each domain, we identify specific extractive mechanisms—technical or institutional habits that centralize control—and contrast them with high-agency pathways where communities are actively reclaiming authority.

Although mechanisms associated with each domain can appear at multiple points in AI development, consistent tendencies emerge across the corpus: decisions about relations often arise upstream as problems are framed; labor arrangements cluster within mid-stream annotation workflows; representational decisions crystallize where ontologies

⁴⁶⁹ and preprocessing pipelines are defined; infrastructural conditions span stages but become most visible as systems scale;
⁴⁷⁰ and governance concerns intensify downstream as models move toward evaluation and deployment. These tendencies
⁴⁷¹ help situate each domain without implying a fixed or linear pipeline.
⁴⁷²

⁴⁷³ Each domain is introduced through a small set of examples that surface the mechanisms we observed across the
⁴⁷⁴ corpus. These examples are intended as points of entry into a broader landscape. The larger set of mappings, summaries,
⁴⁷⁵ and domain categorizations is available in the datasheet for readers who wish to trace these patterns in greater depth.
⁴⁷⁶

⁴⁷⁷ 4.1 Data Relations

⁴⁷⁸ Data relations define the structural terms of engagement between model developers and the communities from whom
⁴⁷⁹ knowledge is derived. In mainstream industry discourse, these engagements are frequently reduced to legalistic questions
⁴⁸⁰ of copyright compliance or static “terms of service.” However, our corpus reveals that these legal frameworks often serve
⁴⁸¹ to obscure the underlying power dynamics [119]. Relations are not merely contractual; they are the primary site where
⁴⁸² agency is either stripped or substantiated. In extractive regimes, relations are characterized by the severance of ties
⁴⁸³ between data and its creators; as Leanne Betasamosake Simpson articulates, “extraction removes all of the relationships
⁴⁸⁴ that give whatever is being extracted meaning” [79]. High-agency relations, conversely, position data production as a
⁴⁸⁵ negotiated partnership where community authority persists even after data is collected.
⁴⁸⁶

⁴⁸⁷ **The assumption of availability** constitutes the primary mechanism of extractive relations. Technical workflows for
⁴⁸⁸ foundation models frequently operate on the premise that any data accessible on the public web is a “standing reserve”
⁴⁸⁹ available for ingestion. This logic converts public existence into implicit consent. Large-scale scraping initiatives, such
⁴⁹⁰ as the corpora used to train models like CLIP [142, 143] or T5 [128], for example, bypass the negotiation of relationship
⁴⁹¹ entirely, including legally, by treating the act of publication as a forfeiture of rights [75, 139]. Relation-less forms of data
⁴⁹² production systematically ignore the contextual intent of the data creator, whether it be repurposing religious texts or
⁴⁹³ intimate narratives as generic linguistic tokens, none of which are “just data” [63]. By removing the requirement to ask,
⁴⁹⁴ the assumption of availability structurally precludes the possibility of refusal, rendering the relationship unilateral.
⁴⁹⁵

⁴⁹⁶ **Transactional asymmetry** reinforces this extraction by decoupling value generation from risk. This manifests in
⁴⁹⁷ “digital extractivism,” where Global Majority communities provide the raw material while the risks—such as the loss of
⁴⁹⁸ privacy or the commodification of cultural heritage—are externalized back to them [67]. The dynamic functions through
⁴⁹⁹ “accumulation by dispossession,” where the terms of engagement are dictated by the extractor, treating communities
⁵⁰⁰ as resources rather than partners [160]. Relational asymmetries are both economic and epistemic. AI developers
⁵⁰¹ gain a model of the world, while communities lose control over how they are represented within it, often leading to
⁵⁰² “opportunity loss” where resources are withheld based on extractive profiling [146].
⁵⁰³

⁵⁰⁴ **High-agency relations** counter these mechanisms by shifting from static terms of service to dynamic and revocable
⁵⁰⁵ consent. Rather than viewing consent as a one-time gatekeeping mechanism, high-agency approaches frame it as an
⁵⁰⁶ ongoing relationship. The Speech Accessibility Project [2] and other initiatives that engage disability communities
⁵⁰⁷ aptly demonstrate how relationships can precede collection: communities are partners who co-define the terms of
⁵⁰⁸ engagement before recruiting paid volunteers and help ensure the protocol aligns with community safety needs
⁵⁰⁹ [2, 117]. Similarly, feminist frameworks for “embodied consent” argue for agreements that are specific, enthusiastic,
⁵¹⁰ and revocable, challenging the broad permissions usually buried in click-through agreements [151].
⁵¹¹

⁵¹² In Indigenous contexts, high-agency relations manifest as relational sovereignty. Te Hiku Media’s approach to Māori
⁵¹³ data rejects the concept of open-source availability in favor of whanaungatanga (connection/relationship), where data
⁵¹⁴ access is determined by the strength and trust of the relationship between the parties [28, 60]. This reintroduces friction
⁵¹⁵
⁵¹⁶ Manuscript submitted to ACM
⁵¹⁷
⁵¹⁸
⁵¹⁹
⁵²⁰

521 into the data pipeline by design: access is not a default state but a negotiated privilege that requires maintaining a
522 relationship with the originating community [81]. By replacing the assumption of availability with permissioned access,
523 these models force a structural acknowledgment of community agency.

525
526
527 **Key takeaway:** Data relations determine the flow of agency. Extractive mechanisms rely on the assumption
528 of availability, treating public data as a resource to be mined and severing the link between creators and their
529 data. This creates transactional asymmetry, where developers capture value while communities bear the risk.
530 High-agency relations replace this with dynamic consent and relational sovereignty, ensuring that data production
531 remains a negotiated partnership where community authority persists throughout the technical lifecycle.
532
533

534 4.2 Data Labor

535 Data labor encompasses the human energy and interpretive judgment required to bridge the gap between raw information
536 and computational capability. While often obscured by the metaphor of “autonomous” AI, our corpus confirms that
537 model performance remains strictly dependent on human workers who select, annotate, validate, and moderate content
538 [26, 54, 85, 104]. In extractive regimes, this labor is characterized by value capture, where the semantic value generated
539 by human judgment is stripped from the worker and concentrated in the model, often leaving the contributor with
540 little to no recognition or economic return. High-agency approaches, conversely, frame labor as expertise, positioning
541 annotators as skilled contributors whose situated knowledge is essential to system quality.

542 **Invisibilization** by design constitutes the primary mechanism of labor extraction. Dataset and platform architectures
543 are frequently designed to present corpora as neutral technical artifacts rather than products of human judgment,
544 masking the interpretive decisions embedded in every labeled example [96]. This structural opacity serves to commodify
545 the worker; by decomposing complex cultural tasks into fragmented “microtasks,” platforms strip the work of its context,
546 rendering the worker interchangeable and the labor invisible [39]. Here, a structural design choice renders the human
547 contribution indistinguishable from the system’s output, with the upshot of systematically preventing workers from
548 asserting authorship claims or contesting the terms of their participation.

549 **Reciprocity failure** reinforces this dynamic by extracting labor without returning value. This manifests most clearly
550 in “unwitting” labor, where user interactions (e.g., solving CAPTCHAs, tagging photos, or correcting autocomplete
551 suggestions) are harvested to train models without the user’s explicit knowledge or compensation [17, 102]. In Global
552 Majority contexts, this mechanism appears in the outsourcing of trauma-inducing content moderation or complex
553 annotation to workers in low-income regions, who perform essential semantic labor for wages that do not reflect the
554 cognitive intensity of the work [163]. The system is optimized to externalize the costs of dataset construction to the
555 worker while centralizing the economic benefits.

556 **High-agency labor** counters these mechanisms by restructuring the economic and attributional relationship
557 between modelers and workers. These approaches restore context and visibility to the labor process. The organization
558 Karya, for example, demonstrates how data collection can function as a tool for economic redistribution; by establishing
559 ethical wage floors and data ownership structures for rural Indian workers, they reframe annotation as a skilled,
560 compensated profession [3]. Similarly, the Masakhane community creates participatory research models where African
561 language speakers function not as passive data subjects, but as credited authors and technical collaborators throughout
562 the pipeline [109]. Emerging initiatives like Ubuntu-AI attempt to encode these rights directly into the data lifecycle

573 through profit-sharing mechanisms, ensuring that artists and creators retain a stake in the value their data generates
574 [108].
575

576 577 **Key takeaway:** Data labor is economic and political. Extractive mechanisms rely on invisibilization by design,
578 decomposing expert judgment into fragmented tasks to obscure the worker and facilitate value capture. This
579 severs the link between labor and downstream value. High-agency approaches replace this with labor as expertise,
580 ensuring that contributions are visible, attributed, and compensated as skilled work that persists within the
581 technical system.
582

583 5 Data Representation

584 Data representation determines how communities become computationally visible. It is less a question of inclusion
585 ratios or diversity statistics than one of epistemic authority: who gets to define the categories, taxonomies, and labels
586 that structure the digital world. In extractive regimes, representation creates visibility without power, often flattening
587 complex, relational identities into rigid categories that facilitate control or consumption. High-agency approaches,
588 conversely, frame representation as plural epistemologies, ensuring that data structures reflect community worldviews
589 rather than forcing local knowledge into universalizing boxes.

590 **Ontological imposition** constitutes the primary mechanism of representational extraction. Institutional problem
591 formulation often imposes external taxonomies on communities before they even enter the pipeline. This manifests as
592 “data universalism,” where Western logics of property and individualism are treated as neutral defaults, overwriting
593 Indigenous ontologies that emphasize relationality and collective stewardship [86, 99]. For example, psychological
594 frameworks developed in “WEIRD” (Western, Educated, Industrialized, Rich, and Democratic) contexts fail to map
595 onto collective ontologies, yet are deployed globally as standard [40, 98, 101]. This mechanism ensures that even when
596 diverse data is collected, it is structurally distorted to fit the model’s worldview, rendering specific cultural meanings
597 “absent” even within inclusion efforts [12].

598 **Context stripping** reinforces this dynamic during annotation and processing. To make data “model-ready,” complex
599 human experiences must be converted into discrete labels. This process often relies on “lazy” data practices that collapse
600 distinct protected attributes like race and ethnicity into coarse categories to satisfy technical constraints, erasing
601 intersectional realities [147]. Annotation workflows that lack community-defined criteria force workers to resolve
602 ambiguity by falling back on institutional defaults, which appear neutral but encode specific cultural biases [138].
603 Automated filtering pipelines compound this by removing content that signals non-normative identities under the guise
604 of “cleaning,” disproportionately purging data from non-Western contexts or disability communities [93].

605 **Synthetic displacement** introduces a new mechanism of extraction: representation without presence. As privacy
606 regulations tighten, developers increasingly turn to synthetic data (e.g., fabricated medical records, artificial faces, and
607 simulated identities) to populate datasets. While this bypasses the need for individual consent [84], it severs the link
608 between representation and reality. Communities become represented in systems they never participated in, inheriting
609 the risks of misidentification or caricature without any pathway to contest how they are depicted [167]. The resulting
610 “diversity-washing” effect is such that models appear inclusive while structurally excluding actual community members.

611 **High-agency representation** counters these mechanisms by building pluralistic and community-grounded corpora.
612 These initiatives prioritize depth and context over scale. For instance, the Abundant Intelligences project reimagines
613 AI development through Indigenous knowledge systems, refusing to separate data from the land and relations that

625 generate it [87]. Similarly, examples from Africa and Oceania demonstrate how regional collaborations can curate
626 datasets that serve local linguistic needs—such as the InkubaLM model—rather than adapting to global benchmarks
627 [43, 162]. By maintaining representational authority, these projects ensure that visibility serves community goals, such
628 as language revitalization, rather than external commodification.

629
630
631
632 **Key takeaway:** Data representation is epistemic and political. Extractive mechanisms rely on ontological
633 imposition and context stripping, imposing external taxonomies and flattening meaning to fit technical defaults.
634 This treats visibility as neutral even when it creates exposure. High-agency approaches replace this with plural
635 epistemologies, grounding representation in community-defined categories and preserving the specificity of local
636 knowledge against universalizing standards.
637
638

640 6 Data Infrastructure

641 Data infrastructure allocates capacity and determines where data lives, who controls access, and how material circulates
642 across model pipelines. While often treated as neutral plumbing designed for efficiency, infrastructure emerges in the
643 literature as a primary site of political contestation. In extractive regimes, infrastructure is configured to maximize
644 velocity and volume, creating technical conditions where consent and context are structurally impossible to maintain.
645 High-agency approaches, conversely, design for traceability and distribution, ensuring that community authority travels
646 with the data.
647

648 **Centralization without governance** configures extraction at an industrial scale. Foundation-model development
649 relies on automated pipelines that ingest content from large-scale web sources such as Common Crawl or LAION to
650 maximize throughput [41, 142, 143]. This configuration privileges actors with substantial compute resources and treats
651 data availability as a default condition. The asymmetry is infrastructural: collection mechanisms operate at speeds that
652 make oversight and contestation structurally unworkable for data subjects [168].
653

654 **Benchmark infrastructures** act as gatekeeping mechanisms that enforce dominant (Western) epistemologies as
655 universal standards. Reliance on a narrow set of legacy datasets, such as ImageNet [38] and MS COCO [90], entrenches
656 specific linguistic, cultural, and demographic assumptions as infrastructural norms [39, 80]. Because creating culturally
657 specific alternatives requires substantial institutional support, Euro-American category systems persist as de facto
658 standards through infrastructural path dependence [80].
659

660 **Provenance compression** serves as the third mechanism, severing datasets from their originating communities
661 and the relational contexts of their creation. Contemporary web-scrape datasets often operate through severe documen-
662 tation gaps—reinforcing “web-as-platform” assumptions that treat public accessibility as permission to extract [139].
663 Infrastructure that treats provenance as optional enables downstream actors to shift responsibility for data quality and
664 rights onto untraceable contributors [91]Longpre et al., 2024).

665 **High-agency infrastructure** counters these mechanisms by embedding community-defined constraints directly
666 into technical architectures. Federated and distributed systems shift authority by enabling collaboration without
667 centralizing data. Emerging frameworks for “data spaces” allow communities to retain local control over storage and
668 access while supporting model development [48, 64]. Similarly, stewardship-based architectures like Masakhane’s
669 distributed research platforms operationalize co-designed metadata standards, ensuring that data does not become
670 “loose” but remains tethered to its community of origin [109].
671

677

678

679

680

681

682

683

684

685

686

687

688

689

690

691 **Key takeaway:** Data infrastructure is about capacity and provenance. Extractive architectures rely on centralization without governance to maximize velocity and provenance compression to sever data from its originating obligations. This makes extraction structurally easy and accountability expensive. High-agency alternatives deploy federated and distributed systems, redistributing capacity to ensure that community authority remains technically enforceable as data circulates.

692

693

694

695

696

7 Data Governance

697

698 Governance establishes the rule-sets that authorize data production: it determines when collection is legitimate, what contextual grounding is required, and who holds authority over circulation. These rules operate upstream of participation, labor, and representation, guiding the conditions under which data production becomes legitimate. High-asymmetry governance frameworks create wide discretionary space for extractive practices, whereas high-agency governance embeds community control directly into the structures that shape data lifecycles.

699

700

701

702

703

704

705

706 **Regulatory arbitrage** constitutes the primary mechanism of extractive governance. Often termed “ethics dumping,” 707 this practice exploits fragmented global regulations to harvest data in regions with weaker protections, converting 708 behavioral interactions into institutional assets without oversight [158]. This dynamic transforms regulatory variation 709 into a resource for extraction: vulnerable populations in low- and middle-income countries may receive limited digital 710 services (like Free Basics) in exchange for extensive, uncompensated data harvesting [110]. Coercive collection in 711 humanitarian settings, such as biometric registration in Ethiopian refugee camps, further illustrates how governance 712 gaps allow institutions to bypass the consent standards required in their home jurisdictions [159].

713

714

715 **Open-loop extraction** reinforces this asymmetry by decoupling deployment from accountability. Models trained 716 on narrow, Western-centric data are frequently deployed globally, shifting the burden of performance failures—such 717 as diagnostic errors in healthcare AI—onto underserved communities [8, 112]. This mechanism externalizes risk: 718 communities excluded from the governance of training data nonetheless become sources of performance feedback 719 during deployment. Their interactions refine the system, yet they possess no authority to challenge the model’s adequacy 720 or recall the data they generate [161]. Governance here functions to protect the model developer’s intellectual property 721 while leaving the data subject’s sovereignty unprotected.

722

723

724

725

726

727

728

729 **High-agency governance** counters these mechanisms through sovereignty-based licensing and critical refusal. 730 Rather than relying on open-access defaults, these approaches encode community authority into the legal terms of the 731 data itself. The Kaitiakitanga License, developed by Te Hiku Media, exemplifies this by legally binding data usage to 732 Māori tikanga (protocols), preventing extractive reuse by third parties [157]. Similarly, the Esethu Framework for African 733 language data establishes sovereignty provisions that mandate community benefit-sharing and protect annotators 734 [129]. Beyond licensing, critical refusal operates as a form of affirmative governance. By setting ex-ante boundaries on 735 participation, communities assert that unreadability is a safety condition. Longstanding tactics of opacity and masking 736 establish practical limits on what institutions may extract [21]. When viewed as governance, refusal is not a lack of 737 data; it is an enforcement of sovereignty that limits extractive reach by design [49].

738

739

740

741

742

743

744

745

746

747

748

749

750

751

752

753

754

755

729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
Key takeaway: Data governance distinguishes accountability from exploitable discretion. Extractive mechanisms rely on regulatory arbitrage (ethics dumping) and open-loop extraction, engaging in collection without contextual grounding and turning deployment into unconsented data acquisition. High-agency approaches establish sovereignty-based licensing and critical refusal, creating enforceable preconditions that align data production with community-defined control and ensure authority persists beyond the point of collection.

8 Discussion

Our analysis of 350 sources across academic and grey literature reveals that AI data production is not merely a logistical preliminary to model development, but a distinct sociotechnical site where power is negotiated, contested, and encoded. By synthesizing evidence across the A/D/C framework, we identify a clear divergence: extractive practices that prioritize scale, opacity, and labor externalization, versus high-agency pathways that prioritize relationality, sovereignty, and context. Noteably, the five analytic domains we identify do not distribute evenly across the ML pipeline. Instead, the sources that comprise each domain cluster around the structural moments where key mechanisms take effect: Data Relations concentrates upstream in problem formulation and data selection; Data Labor anchors mid-pipeline annotation and enrichment; Data Representation spans early- to mid-pipeline ontology and preprocessing; Data Infrastructure forms a cross-cutting substrate most visible in mid-to-downstream development; and Data Governance clusters downstream where deployment, accountability, and sovereignty become salient. This patterned distribution indicates that extractive dynamics are not random but structurally embedded within distinct, yet interrelated, pipeline junctures.

For the HCI community, these findings suggest a critical reframing. While HCI has successfully interrogated downstream AI interaction (how users experience models) and mid-stream model behavior (bias and fairness), the upstream processes of data creation remain undertheorized in design venues. Below, we discuss how HCI scholars and practitioners can operationalize high-agency practices by treating data production as a primary site of design intervention. In doing so, we extend the nascent HCI scholarship that examines data production through the lens of data laborers and data subjects [73, 74, 97, 137, 140, 164].

8.1 Reframing Data Production as “Upstream” Design

Our findings challenge the industry norm of treating data as “found” infrastructure (Era 3). Instead, the evidence suggests data production is a series of design decisions—regarding relations, labor, and representation—that are often irreversible once encoded into a model. This recognition prompts us to argue that the “user” in human-centered AI must expand to include the data contributor—the artist, the annotator, the community member—whose agency is often circumvented by upstream infrastructure. Within HCI, this circumvention has predominantly been investigated in studies of data labor, which reveal how data annotators are frequently reduced to an interchangeable resource thereby constraining their subjectivities and interpretive work [73, 97, 164]. Building on this work, our review makes clear the various design choices—such as crowdsourcing interfaces that atomize tasks to obscure the worker’s context (Data Labor) or scraping pipelines that strip provenance metadata (Data Infrastructure)—that can circumvent agency and enforce extraction by design.

For HCI, this implies that data curation is a form of interaction design. The high-agency pathways our review surfaces make clear that alternative designs are possible. The community-led initiatives—such as Masakhane’s participatory NLP [94] or Māori data sovereignty protocols [81, 105]—succeed not by “fixing” extraction after the fact, but by designing

781 relational friction into the process. They replace the seamless, frictionless extraction of web scraping with protocols
782 that require consent, negotiation, and maintenance.
783

784 8.2 Toward High-Agency Practices: Implications for HCI 785

786 Moving beyond critique, our analysis of high-agency pathways points toward concrete mechanisms for less-extractive
787 AI development. We map these implications to three key shifts for HCI practice.
788

789 8.2.1 *From Universal Representation to Pluralistic “Small Data”*. The dominance of massive, web-scraped corpora (Era
790 3) enforces a universalizing worldview that erases minority contexts and agency in general. Our findings suggest that
791 “de-biasing” these massive datasets is often less effective than building smaller, community-sovereign corpora. Indeed,
792 critiques of contemporary efforts to build more “inclusive” or “de-biased” technologies often highlight a technosolutionist
793 trap whereby the issue is purportedly addressed through large-scale capture of data about a community or culture
794 without meaningful agency [25, 126, 127]. Our review reinforces how failing to allow communities to set the terms of
795 inclusion for their data can inadvertently perpetuate extraction under the guise of inclusion. This extends critiques
796 of “fair” AI that don’t fundamentally shift power [71]. In contrast, high-agency pathways highlighted in our review
797 demonstrate how different actors and groups are proactively responding to extractive data capture by imagining
798 and building alternatives. For instance, efforts by Te Hiku Media to develop community led language datasets and
799 technologies [43, 157] and the Community-driven African Next Voices project [172] directly counter efforts by big
800 tech to seek out and capture cultural knowledge and data, by instead keeping data governed by communities. While
801 authors, organizations, and initiatives offer unique contributions, their coordination and totality points to systemic
802 alternatives that start with community needs and maintain community control, thus creating their own conditions
803 for thriving rather than adapting to external constraints. By developing their own evaluation criteria, publication
804 venues, funding mechanisms, and governance protocols, they establish parallel infrastructures that operate according
805 to different principles: sovereignty rather than extraction, reciprocity rather than accumulation, cultural preservation
806 rather than homogenization and standardization.
807

808 The HCI community has an opportunity to advance high-agency efforts by investing in federated data spaces rather
809 than centralized lakes. We need infrastructure that allows models to learn from community data without that data
810 ever leaving the community’s local storage or jurisdiction (e.g., federated learning tailored for Indigenous sovereignty
811 [Garcia et al., 2024]). Furthermore, valuation metrics in AI research must shift away from scale at the expense of care
812 [59, 69, 137, 171]. We encourage the HCI community to value (and publish) contributions that curate high-context,
813 small-scale datasets with clear governance protocols, rather than rejecting them for lacking the scale of foundation
814 model benchmarks.
815

816 8.2.2 *From Transactional Labor to Relational Provenance*. Our review highlights reciprocity failure and labor invisibility
817 as central extractive patterns. This transactional model commodifies the work of data laborers—including annotators,
818 content creators, and community members—distancing those performing the work from those capturing the value and
819 contributing to the perceived “magic” of AI [135]. This is further complicated by opaque data collection practices that
820 often make the data creator unaware of their contribution. While there is growing interest in data provenance as a
821 key intervention point for mitigating harm of AI technologies [91, 166], our review affirms how the current dominant
822 paradigm of data production is in tension with this end goal. We argue that this tension is, in part, a design challenge
823 for HCI researchers and call upon the community to explore how data capture and sharing platforms might implement
824 provenance-tracking mechanisms by design.
825

833 A provenance-first design approach could involve binding labor and authorship metadata to individual data points
834 so that creators retain "credits" (similar to the Ubuntu-AI model [108]) that persist through the pipeline. More broadly,
835 there is growing recognition that data annotation is fundamentally subjective and interpretive, often shaped by the
836 sociocultural backgrounds and lived experiences of annotators [44? ?]. This motivates the design of data annotation
837 processes and infrastructure that allows workers to signal ambiguity, refuse tasks that violate community norms, and
838 capture disagreements in a structured form rather than forcing a choice that flattens cultural context. By doing so, the
839 HCI community can enable downstream pluralistic modeling approaches that can handle meaningful divergences in
840 perspectives [103].
841

842 *8.2.3 From Open-Loop Extraction to Closed-Loop Governance.* The governance gaps identified in our review show that
843 once data is scraped, communities often lose control. In contrast, community-led governance approaches exemplify an
844 alternative. For example, Māori data sovereignty frameworks in Aotearoa New Zealand demonstrate a coordinated
845 ecosystem: the Māori Data Sovereignty Network develops governance protocols, Te Hiku Media creates community-led,
846 culturally appropriate datasets and benchmarks, and the Kaitiakitanga License embeds community authority into legal
847 frameworks. Such agency-oriented practices require dynamic consent and enforceable boundaries and necessitate
848 technical implementations of Sovereignty-Based Licensing. HCI scholars can develop and standardize machine-readable
849 licenses (similar to Creative Commons but for ML training) that explicitly forbid certain downstream uses (e.g., military
850 application, generative mimicry) and trigger benefit-sharing clauses [129, 157].
851

852 To operationalize this, Institutional Review Boards (IRBs) and conference ethics reviews must look upstream, enforcing
853 data transparency standards that treat data collection as a distinct object of ethical inquiry [10, 52]. Within academic
854 peer review, data production is increasingly within scope of ethical inquiry. However, the focus remains largely on
855 individual privacy and consent within papers presenting novel datasets, rather than deeper inquiries into the conditions
856 under which data is produced and the extent to which communities retain any rights of refusal. This necessitates
857 new review paradigms that prioritize community consent in addition to individual terms of service, echoing calls for
858 power-aware approaches that allow communities to attest and refuse data extraction [49, 71]. By scrutinizing these
859 data cascades at the source [134], the review process can identify where the agency of the data contributor has been
860 circumvented.
861

862 **8.3 Methodological Contributions: The Value of Multivocality**

863 Finally, this paper validates the utility of the Multivocal Literature Review (MLR) for investigating sociotechnical harm.
864 A significant portion of our high-agency evidence came not from peer-reviewed academic venues, but from "grey"
865 literature—community manifestos, tribal resolutions, and worker inquiries. If we had limited our scope to academic
866 "white" literature, we would have successfully diagnosed the harms of extraction (which are well-documented in
867 academia) but missed the existing alternatives (which are often documented in policy and community organizing).
868 For the HCI community, this underscores that "state-of-the-art" knowledge regarding justice and equity often resides
869 outside the academy, as does the general "state-of-practice." Future work on AI harms should adopt multivocal methods
870 to ensure that community-generated resistance and innovation are recognized as rigorous evidence.
871

872 **9 Conclusions**

873 As AI development consolidates around foundation models trained on internet-scale scrapes, the risk of deepening
874 extractive relations is acute. However, this trajectory is not inevitable. By analyzing the data production pipeline
875

885 through the lens of Data Relations, Labor, Representation, Infrastructure, and Governance, we see that every dataset
886 is a record of power relations. This paper contributes a taxonomy of these relations, offering HCI a diagnostic tool
887 to identify extraction and a catalog of precedents for resistance. The shift to less-extractive AI requires more than
888 better algorithms; it requires designing the upstream sociotechnical infrastructures that determine whose knowledge
889 counts, how it is valued, and who governs its future. Our review affirms that a less-extractive future is not merely an
890 aspiration; it is actively being built by communities pursuing alternative to the status quo.
891

892

893

894

895

896

897

898

899

900

901

902

903

904

905

906

907

908

909

910

911

912

913

914

915

916

917

918

919

920

921

922

923

924

925

926

927

928

929

930

931

932

933

934

935

936

937 References

- 938 [1] [n. d.]. Data Workers Inquiry. <https://data-workers.org/>
- 939 [2] [n. d.]. Speech Accessibility Project. <https://speechaccessibilityproject.beckman.illinois.edu>
- 940 [3] Basil Abraham, Danish Goel, Divya Siddarth, Kalika Bali, Manu Chopra, Monojit Choudhury, Pratik Joshi, Preethi Jyoti, Sunayana Sitaram, and Vivek Seshadri. 2020. Crowdsourcing Speech Data for Low-Resource Languages from Low-Income Workers. In *Proceedings of the Twelfth Language Resources and Evaluation Conference*, Nicoletta Calzolari, Frédéric Béchet, Philippe Blache, Khalid Choukri, Christopher Cieri, Thierry Declerck, Sara Goggi, Hitoshi Isahara, Bente Maegaard, Joseph Mariani, Hélène Mazo, Asuncion Moreno, Jan Odijk, and Stelios Piperidis (Eds.). European Language Resources Association, Marseille, France, 2819–2826. <https://aclanthology.org/2020.lrec-1.343/>
- 941 [4] Mark Adley, Hayley Alderson, Katherine Jackson, William McGovern, Liam Spencer, Michelle Addison, and Amy O'Donnell. 2024. Ethical and practical considerations for including marginalised groups in quantitative survey research. *International Journal of Social Research Methodology* 27, 5 (2024), 559–574. doi:10.1080/13645579.2023.2228600
- 942 [5] Leah Hope Ajmani, Jasmine C. Foriest, Jordan Taylor, Kyle Pittman, Sarah Gilbert, and Michael Ann DeVito. 2024. Whose Knowledge is Valued? Epistemic Injustice in CSCW Applications. *Proc. ACM Hum.-Comput. Interact.* 8, CSCW2 (Nov. 2024), 523:1–523:28. doi:10.1145/3687062
- 943 [6] Hilary Arksey and Lisa O'Malley. 2005. Scoping studies: towards a methodological framework. *International Journal of Social Research Methodology* 8, 1 (Feb. 2005), 19–32. doi:10.1080/1364557032000119616
- 944 [7] A. Arora, M. Barrett, E. Lee, E. Oborn, and K. Prince. 2023. Risk and the future of AI: Algorithmic bias, data colonialism, and marginalization. *Information and Organization* 33, 3 (2023), 100478. doi:10.1016/j.infoandorg.2023.100478
- 945 [8] Mercy Asiedu, Awa Dieng, Iskandar Haykel, Negar Rostamzadeh, Stephen Pfohl, Chirag Nagpal, Maria Nagawa, Abigail Oppong, Sanmi Koyejo, and Katherine Heller. 2024. The Case for Globalizing Fairness: A Mixed Methods Study on Colonialism, AI, and Health in Africa. arXiv:2403.03357 (March 2024). doi:10.48550/arXiv.2403.03357 arXiv:2403.03357 [cs].
- 946 [9] Marie Battiste. 2005. Indigenous Knowledge: Foundations for First Nations. *Worm Indigenous Nations Higher Education Consortium Journal* (Jan. 2005). https://www.researchgate.net/publication/241822370_Indigenous_Knowledge_Foundations_for_First_Nations
- 947 [10] Emily M. Bender and Batya Friedman. 2018. Data Statements for Natural Language Processing: Toward Mitigating System Bias and Enabling Better Science. *Transactions of the Association for Computational Linguistics* 6 (Dec. 2018), 587–604. doi:10.1162/tacl_a_00041
- 948 [11] Eshta Bhardwaj, Harshit Gujral, Siyi Wu, Ciara Zogheib, Tegan Maharaj, and Christoph Becker. 2024. Machine learning data practices through a data curation lens: An evaluation framework. In *The 2024 ACM Conference on Fairness, Accountability, and Transparency*. ACM, Rio de Janeiro Brazil, 1055–1067. doi:10.1145/3630106.3658955
- 949 [12] Steven Bird. 2024. Must NLP be Extractive? https://drive.google.com/file/d/1hvF7_WQrou6CWZydhyMFTYHnd3ZljV/view?usp=embed_facebook
- 950 [13] Abeba Birhane. 2020. Algorithmic Colonization of Africa. *SCRIPTed* 17, 2 (Aug. 2020), 389–409. doi:10.2966/scrip.170220.389
- 951 [14] Abeba Birhane. 2021. Algorithmic injustice: a relational ethics approach. *Patterns* 2, 2 (Feb. 2021), 100205. doi:10.1016/j.patter.2021.100205
- 952 [15] Abeba Birhane, William Isaac, Vinodkumar Prabhakaran, Mark Diaz, Madeleine Clare Elish, Jason Gabriel, and Shakir Mohamed. 2022. Power to the People? Opportunities and Challenges for Participatory AI. In *Equity and Access in Algorithms, Mechanisms, and Optimization*. ACM, Arlington VA USA, 1–8. doi:10.1145/3551624.3555290
- 953 [16] Abeba Birhane, Elayne Ruane, Thomas Laurent, Matthew S. Brown, Johnathan Flowers, Anthony Ventresque, and Christopher L. Dancy. 2022. The Forgotten Margins of AI Ethics. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency (FAccT '22)*. Association for Computing Machinery, New York, NY, USA, 948–958. doi:10.1145/3531146.3533157
- 954 [17] Briony Blackmore, Michelle Thorp, Andrew Tzer-Yeu Chen, Fabio Morreale, Brent Burmester, Elham Bahmanteymouri, and Matt Bartlett. 2023. Hidden humans: exploring perceptions of user-work and training artificial intelligence in Aotearoa New Zealand. *Kōtuitui: New Zealand Journal of Social Sciences Online* 18, 4 (Oct. 2023), 443–456. doi:10.1080/1177083X.2023.2212736
- 955 [18] Su Lin Blodgett, Solon Barocas, Hal Daumé III, and Hanna Wallach. 2020. Language (Technology) is Power: A Critical Survey of “Bias” in NLP. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, Dan Jurafsky, Joyce Chai, Natalie Schluter, and Joel Tetreault (Eds.). Association for Computational Linguistics, Online, 5454–5476. doi:10.18653/v1/2020.acl-main.485
- 956 [19] John R. Bowman. 1989. *Capitalist Collective Action: Competition, Cooperation and Conflict in the Coal Industry*. Cambridge University Press. Google-Books-ID: fnl6sAYpRLYC.
- 957 [20] Paul T. Brown, Daniel Wilson, Kiri West, Kirita-Rose Escott, Kiya Basabas, Ben Ritchie, Danielle Lucas, Ivy Taia, Natalie Kusabs, and Te Taka Keegan. 2024. Māori Algorithmic Sovereignty: Idea, Principles, and Use. *Data Science Journal* 23, 1 (April 2024). doi:10.5334/dsj-2024-015
- 958 [21] Simone Browne. 2015. *Dark Matters: On the Surveillance of Blackness*. Duke University Press. doi:10.1215/9780822375302
- 959 [22] Jenna Burrell. 2016. How the machine ‘thinks’: Understanding opacity in machine learning algorithms. *Big Data & Society* 3, 1 (2016), 2053951715622512. doi:10.1177/2053951715622512
- 960 [23] Judith Butler. 1990. *Gender Trouble: Feminism and the Subversion of Identity*. Routledge. Google-Books-ID: kuztAAAAMAAJ.
- 961 [24] Stephanie Russo Carroll, Pyrou Chung, Robyn K. Rowe, Susanna Siri, and Walter. 2024. Indigenous Data Sovereignty and the State of Open Data. <https://www.d4d.net/news/indigenous-data-sovereignty-and-the-state-of-open-data>
- 962 [25] Alan Chan, Chinasa T Okolo, Zachary Termer, and Angelina Wang. 2021. The limits of global inclusion in AI development. *arXiv preprint arXiv:2102.01265* (2021).

- [989] [26] Srravya Chandhiramowuli, Alex S. Taylor, Sara Heitlinger, and Ding Wang. 2024. Making Data Work Count. *Proc. ACM Hum.-Comput. Interact.* 8, CSCW1 (April 2024), 90:1–90:26. doi:10.1145/3637367
- [990] [27] Ishita Chordia, Leya Breanna Baltaxe-Admony, Ashley Boone, Alyssa Sheehan, Lynn Dombrowski, Christopher A Le Dantec, Kathryn E. Ringland, and Angela D. R. Smith. 2024. Social Justice in HCI: A Systematic Literature Review. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems (CHI '24)*. Association for Computing Machinery, New York, NY, USA, 1–33. doi:10.1145/3613904.3642704
- [991] [28] Donavyn Coffey. 2021. Māori are trying to save their language from Big Tech. *Wired* (April 2021). <https://www.wired.com/story/maori-language-tech/>
- [992] [29] Cathy J. Cohen. 1997. Punks, Bulldaggers, and Welfare Queens: The Radical Potential of Queer Politics? *GLQ: A Journal of Lesbian and Gay Studies* 3, 4 (1997), 437–465.
- [993] [30] Combahee River Collective. 1977. (1977) The Combahee River Collective Statement •. <https://www.blackpast.org/african-american-history/combahee-river-collective-statement-1977/>
- [994] [31] Patricia Hill Collins. 2002. *Black Feminist Thought: Knowledge, Consciousness, and the Politics of Empowerment* (2 ed.). Routledge, New York. doi:10.4324/9780203900055
- [995] [32] Ned Cooper, Tiffanie Horne, Gillian R Hayes, Courtney Heldreth, Michal Lahav, Jess Holbrook, and Lauren Wilcox. 2022. A Systematic Review and Thematic Analysis of Community-Collaborative Approaches to Computing Research. In *CHI Conference on Human Factors in Computing Systems*. ACM, New Orleans LA USA, 1–18. doi:10.1145/3491102.3517716
- [996] [33] Matthew Cotton. 2017. Fair fracking? Ethics and environmental justice in United Kingdom shale gas policy and planning. *Local Environment* 22, 2 (Feb. 2017), 185–202. doi:10.1080/13549839.2016.1186613
- [997] [34] Payton Croskey, Fabian Offert, Jennifer Jacobs, and Kai M. Thaler. 2025. Liberatory Collections and Ethical AI: Reimagining AI Development from Black Community Archives and Datasets. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency (FAccT '25)*. Association for Computing Machinery, New York, NY, USA, 900–913. doi:10.1145/3715275.3732058
- [998] [35] Íñigo de Troya, Jacqueline Kernahan, Neelke Doorn, Virginia Dignum, and Roel Dobbe. 2025. Misabstraction in Sociotechnical Systems. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency (FAccT '25)*. Association for Computing Machinery, New York, NY, USA, 1829–1842. doi:10.1145/3715275.3732122
- [999] [36] Jeffrey Dean. 2019. The Deep Learning Revolution and Its Implications for Computer Architecture and Chip Design. arXiv:1911.05289 (Nov. 2019). doi:10.48550/arXiv.1911.05289 arXiv:1911.05289 [cs].
- [1000] [37] Vine Deloria and Clifford M. Lytle. 1998. *The Nations Within: The Past and Future of American Indian Sovereignty*. University of Texas Press. Google-Books-ID: FLgEf5kGLWQC.
- [1001] [38] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. 2009. ImageNet: A large-scale hierarchical image database. In *2009 IEEE Conference on Computer Vision and Pattern Recognition*. IEEE, Miami, FL, 248–255. doi:10.1109/CVPR.2009.5206848
- [1002] [39] Emily Denton, Alex Hanna, Razvan Amironesei, Andrew Smart, and Hilary Nicole. 2021. On the genealogy of machine learning datasets: A critical history of ImageNet. *Big Data & Society* 8, 2 (2021), 2053951721103595. doi:10.1177/2053951721103595
- [1003] [40] Lindsey DeWitt Prat, Olivia Nercy Ndlovu Lucas, Christopher Golias, and Mia Lewis. 2024. Decolonizing LLMs: An Ethnographic Framework for AI in African Contexts. *EPIC Proceedings* (2024), 45–84. <https://doi.org/10.1111/epic.12196>
- [1004] [41] Jesse Dodge, Maarten Sap, Ana Marasović, William Agnew, Gabriel Ilharco, Dirk Groeneveld, Margaret Mitchell, and Matt Gardner. 2021. Documenting Large Webtext Corpora: A Case Study on the Colossal Clean Crawled Corpus. arXiv:2104.08758 (2021). doi:10.48550/arXiv.2104.08758 arXiv:2104.08758 [cs].
- [1005] [42] Pedro Domingos. 2012. A few useful things to know about machine learning. *Commun. ACM* 55, 10 (Oct. 2012), 78–87. doi:10.1145/2347736.2347755
- [1006] [43] Suzanne Duncan, Gianna Leoni, Lee Steven, Keoni Mahelona, and Peter-Lucas Jones. 2024. Fit for our purpose, not yours: Benchmark for a low-resource, Indigenous language. <https://openreview.net/forum?id=w5jfjyvsRq3#discussion>
- [1007] [44] Mark Diaz, Ian Kivlichan, Rachel Rosen, Dylan Baker, Razvan Amironesei, Vinodkumar Prabhakaran, and Remi Denton. 2022. CrowdWorkSheets: Accounting for Individual and Collective Identities Underlying Crowdsourced Dataset Annotation. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency (FAccT '22)*. Association for Computing Machinery, New York, NY, USA, 2342–2351. doi:10.1145/3531146.3534647
- [1008] [45] Pedro Ferreira. 2024. Examining the “Local” in ICT4D: A Postcolonial Perspective on Participation. In *Proceedings of the CHI Conference on Human Factors in Computing Systems*. ACM, Honolulu HI USA, 1–13. doi:10.1145/3613904.3642748
- [1009] [46] Raymond Fok, Alexa Siu, and Daniel S. Weld. 2025. Toward Living Narrative Reviews: An Empirical Study of the Processes and Challenges in Updating Survey Articles in Computing Research. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems (CHI '25)*. Association for Computing Machinery, New York, NY, USA, 1–10. doi:10.1145/3706598.3714047
- [1010] [47] Miranda Fricker. 2007. *Epistemic injustice: power and the ethics of knowing*. Oxford university press, Oxford.
- [1011] [48] Ana García, Savvas Rogotis, Eimear Farrell, Tobias Guggenberger, Arash Hajikhani, Atte Kinnula, Marko Komssi, and Tuomo Tuikka. 2024. Generative AI and Data Spaces: White Paper. (2024).
- [1012] [49] Patricia Garcia, Tonia Sutherland, Niloufar Salehi, Marika Cifor, and Anubha Singh. 2022. No! Re-imagining Data Practices Through the Lens of Critical Refusal. *Proc. ACM Hum.-Comput. Interact.* 6, CSCW2 (Nov. 2022), 315:1–315:20. doi:10.1145/3557997
- [1013] [50] Vahid Garousi, Michael Felderer, and Mika V. Mäntylä. 2016. The need for multivocal literature reviews in software engineering: complementing systematic literature reviews with grey literature. In *Proceedings of the 20th International Conference on Evaluation and Assessment in Software*

- 1041 Engineering (EASE '16). Association for Computing Machinery, New York, NY, USA, 1–6. doi:10.1145/2915970.2916008
- 1042 [51] Vahid Garousi, Michael Felderer, and Mika V. Mäntylä. 2019. Guidelines for including grey literature and conducting multivocal literature reviews
1043 in software engineering. *Information and Software Technology* 106 (Feb. 2019), 101–121. doi:10.1016/j.infsof.2018.09.006
- 1044 [52] Timnit Gebru, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daumé III, and Kate Crawford. 2021.
1045 Datasheets for datasets. *Commun. ACM* 64, 12 (Nov. 2021), 86–92. doi:10.1145/3458723
- 1046 [53] Katelyn Godin, Jackie Stapleton, Sharon I. Kirkpatrick, Rhona M. Hanning, and Scott T. Leatherdale. 2015. Applying systematic review search
1047 methods to the grey literature: a case study examining guidelines for school-based breakfast programs in Canada. *Systematic Reviews* 4 (Oct. 2015),
1048 138. doi:10.1186/s13643-015-0125-0
- 1049 [54] Mary L Gray and Siddharth Suri. 2019. *Ghost Work: How to Stop Silicon Valley from Building a New Global Underclass*. Houghton Mifflin Harcourt.
- 1050 [55] Michael Haenlein and Andreas Kaplan. 2019. A Brief History of Artificial Intelligence: On the Past, Present, and Future of Artificial Intelligence.
1051 *California Management Review* 61, 4 (Aug. 2019), 5–14. doi:10.1177/0008125619864925
- 1052 [56] Alon Halevy, Peter Norvig, and Fernando Pereira. 2009. The Unreasonable Effectiveness of Data. *IEEE Intelligent Systems* 24, 2 (March 2009), 8–12.
1053 doi:10.1109/MIS.2009.36
- 1054 [57] Siobhan Mackenzie Hall, Samantha Dalal, Raesetje Sefala, Foutse Yuehgooh, Aisha Alaagib, Imane Hamzaoui, Shu Ishida, Jabez Magomere, Lauren
1055 Crais, Aya Salama, and Tejumade Afonja. 2025. The Human Labour of Data Work: Capturing Cultural Diversity through World Wide Dishes.
1056 arXiv:2502.05961 (Feb. 2025). doi:10.48550/arXiv.2502.05961 arXiv:2502.05961 [cs].
- 1057 [58] Jamie Hancock, Sarada Mahesh, Jennifer Cobbe, Jatinder Singh, and Anjali Mazumder. 2024. The tensions of data sharing for human rights:
1058 A modern slavery case study. In *The 2024 ACM Conference on Fairness, Accountability, and Transparency*. ACM, Rio de Janeiro Brazil, 974–987.
1059 doi:10.1145/3630106.3658949
- 1060 [59] Alex Hanna and Tina M. Park. 2020. Against Scale: Provocations and Resistances to Scale Thinking. arXiv:2010.08850 (Nov. 2020). doi:10.48550/
1061 arXiv.2010.08850 arXiv:2010.08850 [cs].
- 1062 [60] Karen Hao. 2022. A new vision of artificial intelligence for the people. *MIT Technology Review* (April 2022). <https://www.technologyreview.com/2022/04/22/1050394/artificial-intelligence-for-the-people/>
- 1063 [61] Karen Hao and Andrea Paola Hernández. 2022. How the AI industry profits from catastrophe. *MIT Technology Review* (April 2022). <https://www.technologyreview.com/2022/04/20/1050392/ai-industry-appen-scale-data-labels/>
- 1064 [62] Donna Haraway. 1988. Situated Knowledges: The Science Question in Feminism and the Privilege of Partial Perspective. *Feminist Studies* 14, 3
1065 (1988), 575–599. doi:10.2307/3178066
- 1066 [63] Ben Hutchinson. 2024. Modeling the Sacred: Considerations when Using Religious Texts in Natural Language Processing. arXiv:2404.14740 (2024).
1067 doi:10.48550/arXiv.2404.14740 arXiv:2404.14740 [cs].
- 1068 [64] Andreas Hutterer and Barbara Krumay. 2024. The adoption of data spaces: Drivers toward federated data sharing. doi:10.24251/HICSS.2024.542
- 1069 [65] Paul Agu Igwe, Nnamdi O. Madichie, and David Gamarie Rugara. 2022. Decolonising research approaches towards non-extractive research.
1070 *Qualitative Market Research: An International Journal* 25, 4 (Jan. 2022), 453–468. doi:10.1108/QMR-11-2021-0135
- 1071 [66] Lilly Irani, Janet Vertesi, Paul Dourish, Kavita Philip, and Rebecca E. Griner. 2010. Postcolonial computing: a lens on design and development. In
1072 *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. ACM, Atlanta Georgia USA, 1311–1320. doi:10.1145/1753326.1753522
- 1073 [67] Neema Iyer, Garnett Achienye, Favour Borokini, Uri Ludger, Neema Iyer, Yahya Syabani, and Yahya Syabani. 2021. Automated Imperialism,
1074 Expansionist Dreams: Exploring Digital Extractivism in Africa. (2021). <https://archive.policy.org/digitalextractivism/>
- 1075 [68] C. L. R. James. 1989. *The black Jacobins: Toussaint l'Overture and the San Domingo revolution* (2. ed., rev ed.). Vintage Books, a Division of Random
1076 House, Inc, New York.
- 1077 [69] Eun Seo Jo and Timnit Gebru. 2020. Lessons from archives: strategies for collecting sociocultural data in machine learning. In *Proceedings of the
1078 2020 Conference on Fairness, Accountability, and Transparency (FAT* '20)*. Association for Computing Machinery, New York, NY, USA, 306–316.
1079 doi:10.1145/3351095.3372829
- 1080 [70] James H Jones. 2008. The Tuskegee syphilis experiment. *The Oxford textbook of clinical research ethics* (2008), 86–96.
- 1081 [71] Pratyusha Kalluri. 2020. Don't ask if artificial intelligence is good or fair, ask how it shifts power. *Nature* 583, 7815 (2020), 169–169.
- 1082 [72] Fernando Kamei, Igor Wiese, Gustavo Pinto, Waldemar Ferreira, Márcio Ribeiro, Renata Souza, and Sérgio Soares. 2022. Assessing the Credibility
1083 of Grey Literature: A Study with Brazilian Software Engineering Researchers. *Journal of Software Engineering Research and Development* 10 (June
1084 2022). doi:10.5753/jserd.2022.1897
- 1085 [73] Shivani Kapania, Alex S Taylor, and Ding Wang. 2023. A hunt for the Snark: Annotator Diversity in Data Practices. In *Proceedings of the 2023 CHI
1086 Conference on Human Factors in Computing Systems*. ACM, Hamburg Germany, 1–15. doi:10.1145/3544548.3580645
- 1087 [74] Reishiro Kawakami and Sukrit Venkatagiri. 2024. The Impact of Generative AI on Artists. In *Proceedings of the 16th Conference on Creativity &
1088 Cognition (CampC '24)*. Association for Computing Machinery, New York, NY, USA, 79–82. doi:10.1145/3635636.3664263
- 1089 [75] Mehtab Khan and Alex Hanna. 2022. The Subjects and Stages of AI Dataset Development: A Framework for Dataset Accountability. *Forthcoming
1090 19 Ohio St. Tech. L.J. (2023)* (2022). doi:10.2139/ssrn.4217148
- 1091 [76] Rob Kitchin, Juliette Davret, Carla M Kayanan, and Samuel Mutter. 2025. Assemblage theory, data systems and data ecosystems: The data
1092 assemblages of the Irish planning system. *Big Data & Society* 12, 3 (2025), 20539517251352822. doi:10.1177/20539517251352822
- 1093 [77] Rob Kitchin and Tracey Lauriault. 2014. Towards Critical Data Studies: Charting and Unpacking Data Assemblages and Their Work. (2014).
1094 <https://papers.ssrn.com/abstract=2474112>

- [78] Lauren Klein and Catherine D'Ignazio. 2024. Data Feminism for AI. In *The 2024 ACM Conference on Fairness, Accountability, and Transparency*. ACM, Rio de Janeiro Brazil, 100–112. doi:10.1145/3630106.3658543
- [79] Naomi Klein. 2013. Naomi Klein Chats with Leanne Simpson about Idle No More. <https://www.yesmagazine.org/social-justice/2013/03/06/dancing-the-world-into-being-a-conversation-with-idle-no-more-leanne-simpson>
- [80] Bernard Koch, Emily Denton, Alex Hanna, and Jacob G. Foster. 2021. Reduced, Reused and Recycled: The Life of a Dataset in Machine Learning Research. arXiv:2112.01716 (Dec. 2021). doi:10.48550/arXiv.2112.01716 arXiv:2112.01716 [cs, stat].
- [81] Tahu Kukutai and Donna Cormack. 2020. *"Pushing the space": Data sovereignty and self-determination in Aotearoa NZ* (1st edition ed.). Routledge, 21–35. <https://www.taylorfrancis.com/reader/read-online/6abf9fc2-820b-4950-b310-eef574873fb/chapter/pdf?context=ubx>
- [82] Hugo Laurençón, Lucile Saulnier, Thomas Wang, Christopher Akiki, Albert Villanova del Moral, Teven Le Scao, Leandro Von Werra, Chenghao Mou, Eduardo González Ponferrada, Huu Nguyen, Jörg Frohberg, Mario Šaško, Quentin Lhoest, Angelina McMillan-Major, Gerard Dupont, Stella Biderman, Anna Rogers, Loubna Ben allal, Francesco De Toni, Giada Pistilli, Olivier Nguyen, Somaieh Nikpoor, Maraim Masoud, Pierre Colombo, Javier de la Rosa, Paulo Villegas, Tristan Thrush, Shayne Longpre, Sebastian Nagel, Leon Weber, Manuel Muñoz, Jian Zhu, Daniel Van Strien, Zaid Alyafeai, Khalid Almubarak, Minh Chien Vu, Itziar Gonzalez-Dios, Aitor Soroa, Kyle Lo, Manan Dey, Pedro Ortiz Suarez, Aaron Gokaslan, Shamik Bose, David Adelani, Long Phan, Hieu Tran, Ian Yu, Suhas Pai, Jenny Chim, Violette Lepercq, Suzana Ilic, Margaret Mitchell, Sasha Alexandra Luccioni, and Yacine Jernite. 2023. The BigScience ROOTS Corpus: A 1.6TB Composite Multilingual Dataset. arXiv:2303.03915 (March 2023). doi:10.48550/arXiv.2303.03915 arXiv:2303.03915 [cs].
- [83] Y. LeCun. [n.d.]. THE MNIST DATABASE of handwritten digits. <http://yann.lecun.com/exdb/mnist/> ([n.d.]). <https://cir.nii.ac.jp/crid/1571417126193283840>
- [84] Peter Lee. 2024. Synthetic Data and the Future of AI. 4722162 (Feb. 2024). <https://papers.ssrn.com/abstract=4722162>
- [85] Tuukka Lehtiöniemi and Minna Ruckenstein. 2022. *Prisoners training AI: Ghosts, Humans and Values in Data Labour*. Routledge, Abingdon, 184–196. doi:10.4324/9781003170884-4
- [86] Jason Edward Lewis, Angie Abdilla, Noelani Arista, Kaipulaumakanolono Baker, Scott Benesiinaabandan, Michelle Brown, Melanie Cheung, Meredith Coleman, Ashley Cordes, Joel Davison, Kūpono Duncan, Sergio Garzon, D. Fox Harrell, Peter-Lucas Jones, Kekuhi Kealiikanakaoole-haililani, Megan Kelleher, Suzanne Kite, Olin Lagon, Jason Leigh, Maroussia Levesque, Keoni Mahelona, Caleb Moses, Isaac ('Ika'aka) Nahuewai, Kari Noe, Danielle Olson, 'Oīwi Parker Jones, Caroline Running Wolf, Michael Running Wolf, Marlee Silva, Skawennati Fragnito, and Hēmi Whaanga. 2020. *Indigenous Protocol and Artificial Intelligence Position Paper*. Indigenous Protocol and Artificial Intelligence Working Group and the Canadian Institute for Advanced Research, Honolulu, HI. doi:10.11573/spectrum.library.concordia.ca.00986506
- [87] Jason Edward Lewis, Hēmi Whaanga, and Ceyda Yolgörmez. 2025. Abundant intelligences: placing AI within Indigenous knowledge frameworks. *AI & SOCIETY* 40, 4 (April 2025), 2141–2157. doi:10.1007/s00146-024-02099-4
- [88] Calvin A. Liang, Sean A. Munson, and Julie A. Kientz. 2021. Embracing Four Tensions in Human-Computer Interaction Research with Marginalized People. *ACM Trans. Comput.-Hum. Interact.* 28, 2 (April 2021), 14:1–14:47. doi:10.1145/3443686
- [89] Andreas Liesenfeld and Mark Dingemanse. 2024. Rethinking open source generative AI: open washing and the EU AI Act. In *The 2024 ACM Conference on Fairness, Accountability, and Transparency*. ACM, Rio de Janeiro Brazil, 1774–1787. doi:10.1145/3630106.3659005
- [90] Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C Lawrence Zitnick. 2014. Microsoft coco: Common objects in context. In *European conference on computer vision*. Springer, 740–755.
- [91] Shayne Longpre, Nikhil Singh, Manuel Cherep, Kushagra Tiwary, Joanna Materzynska, William Brannon, Robert Mahari, Manan Dey, Mohammed Hamdy, Nayan Saxena, Ahmad Mustafa Anis, Emad A. Alghamdi, Vu Minh Chien, Naana Obeng-Marnu, Da Yin, Kun Qian, Yizhi Li, Minnie Liang, An Dinh, Shrestha Mohanty, Deividas Mataciunas, Tobin South, Jianguo Zhang, Ariel N. Lee, Campbell S. Lund, Christopher Klamm, Damien Sileo, Diganta Misra, Enrico Shippole, Kevin Klyman, Lester JV Miranda, Niklas Muennighoff, Seonghyeon Ye, Seungone Kim, Vipul Gupta, Vivek Sharma, Xuhui Zhou, Caiming Xiong, Luis Villa, Stella Biderman, Alex Pentland, Sara Hooker, and Jad Kabbara. 2024. Bridging the Data Provenance Gap Across Text, Speech and Video. arXiv:2412.17847 (Dec. 2024). doi:10.48550/arXiv.2412.17847 arXiv:2412.17847 [cs].
- [92] Qinghua Lu, Liming Zhu, Xiwei Xu, Jon Whittle, Didar Zowghi, and Aurelie Jacquet. 2024. Responsible AI Pattern Catalogue: A Collection of Best Practices for AI Governance and Engineering. *Comput. Surveys* 56, 7 (2024), 1–35. doi:10.1145/3626234
- [93] Kelly Avery Mack, Rida Qadri, Remi Denton, Shaun K. Kane, and Cynthia L. Bennett. 2024. "They only care to show us the wheelchair": disability representation in text-to-image AI models. In *Proceedings of the CHI Conference on Human Factors in Computing Systems (CHI '24)*. Association for Computing Machinery, New York, NY, USA, 1–23. doi:10.1145/3613904.3642166
- [94] Vukosi Marivate. 2021. *Why African natural language processing now? A view from South Africa AfricaNLP*. Mapungubwe Institute for Strategic Reflection (MISTRAL), 126–152. doi:10.2307/jj.12406168.11
- [95] Donald Jr. Martin. 2020. Upgrading the Product Development Process to Foster Machine Learning Fairness and Ethical AI. <https://www.youtube.com/watch?v=1Uyc9SPeYKA>
- [96] Milagros Miceli and Julian Posada. 2022. The Data-Production Dispositif. *Proceedings of the ACM on Human-Computer Interaction* 6, CSCW2 (Nov. 2022), 1–37. doi:10.1145/3555561
- [97] Milagros Miceli, Martin Schuessler, and Tianling Yang. 2020. Between Subjectivity and Imposition: Power Dynamics in Data Annotation for Computer Vision. *Proceedings of the ACM on Human-Computer Interaction* 4, CSCW2 (Oct. 2020), 1–25. doi:10.1145/3415186
- [98] Rada Mihalcea, Oana Ignat, Longju Bai, Angana Borah, Luis Chiruzzo, Zhiping Jin, Claude Kwizera, Joan Nwatu, Soujanya Poria, and Thamar Solorio. 2025. Why AI is WEIRD and shouldn't be this way: towards AI for everyone, with everyone, by everyone. In *Proceedings of the Thirty-Ninth Annual Meeting of the North American Chapter of the Association for Computational Linguistics (NAACL-HLT '25)*, 1–16. doi:10.1145/3555561.3555561

- 1145 *AAAI Conference on Artificial Intelligence and Thirty-Seventh Conference on Innovative Applications of Artificial Intelligence and Fifteenth Symposium*
 1146 *on Educational Advances in Artificial Intelligence (AAAI'25/IAAI'25/EAAI'25, Vol. 39)*. AAAI Press, 28657–28670. doi:10.1609/aaai.v39i27.35092
- 1147 [99] Stefania Milan and Emiliano Treré. 2019. Big Data from the South(s): Beyond Data Universalism. *Television New Media* 20, 4 (May 2019), 319–335.
 doi:10.1177/1527476419837739
- 1148 [100] Shakir Mohamed, Marie-Therese Png, and William Isaac. 2020. Decolonial AI: Decolonial Theory as Sociotechnical Foresight in Artificial
 1149 Intelligence. *Philosophy Technology* 33, 4 (Dec. 2020), 659–684. doi:10.1007/s13347-020-00405-8
- 1150 [101] Cristina Jayme Montiel and Joshua Uyheng. 2022. Foundations for a decolonial big data psychology. *Journal of Social Issues* 78, 2 (2022), 278–297.
 doi:10.1111/josi.12439
- 1151 [102] Fabio Morreale, Elham Bahmaneymouri, Brent Burmester, Andrew Chen, and Michelle Thorp. 2023. The unwitting labourer: extracting humanness
 1152 in AI training. *AI SOCIETY* (May 2023). doi:10.1007/s00146-023-01692-3
- 1153 [103] Aida Mostafazadeh Davani, Mark Diaz, and Vinodkumar Prabhakaran. 2022. Dealing with Disagreements: Looking Beyond the Majority Vote in
 1154 Subjective Annotations. *Transactions of the Association for Computational Linguistics* 10 (2022), 92–110. doi:10.1162/tacl_a_00449
- 1155 [104] James Muldoon, Callum Cant, Boxi Wu, and Mark Graham. 2024. A Typology of AI Data Work. *Big Data and Society* 11, 11 (March 2024).
 doi:10.1177/20539517241232632
- 1156 [105] Luke Munn. 2024. The five tests: designing and evaluating AI according to indigenous Māori principles. *AI SOCIETY* 39, 4 (Aug. 2024), 1673–1681.
 doi:10.1007/s00146-023-01636-x
- 1157 [106] Zachary Munn, Micah D. J. Peters, Cindy Stern, Catalin Tufanaru, Alexa McArthur, and Edoardo Aromataris. 2018. Systematic review or scoping
 1158 review? Guidance for authors when choosing between a systematic or scoping review approach. *BMC Medical Research Methodology* 18, 1 (Nov.
 1159 2018), 143. doi:10.1186/s12874-018-0611-x
- 1160 [107] Diego I. Murguía and Kathrin Böhling. 2013. Sustainability reporting on large-scale mining conflicts: the case of Bajo de la Alumbra, Argentina.
 1161 *Journal of Cleaner Production* 41 (Feb. 2013), 202–209. doi:10.1016/j.jclepro.2012.10.012
- 1162 [108] M. Nayebare, R. Eglash, U. Kimanuku, R. Baguma, J. Mounsey, and C. Maina. 2023. *Interim Report for Ubuntu-AI: A Bottom-up Approach to More
 1163 Democratic and Equitable Training and Outcomes for Machine Learning*. San Francisco. <https://generativejustice.org/uai/>
- 1164 [109] Wilhelmina Nekoto, Vukosi Marivate, Tshinondiwa Matsila, Timi Fasubaa, Taiwo Fagbogungbe, Solomon Oluwole Akinola, Shamsuddeen Muham-
 1165 mad, Salomon Kabongo Kabenamualu, Salomey Osei, Freshia Sackey, Rubungo Andre Niyongabo, Ricky Macharm, Perez Ogayo, Orevaoghene
 1166 Ahia, Musie Meressa Berhe, Mofetoluwa Adeyemi, Masabata Mokgesi-Selinga, Lawrence Okegbemi, Laura Martinus, Kolawole Tajudeen, Kevin
 1167 Degila, Kelechi Ogueji, Kathleen Siminyu, Julia Kreutzer, Jason Webster, Jamil Toure Ali, Jade Abbott, Iroro Orife, Ignatius Ezeani, Idris Abdulkadir
 1168 Dangana, Herman Kamper, Hady Elsahar, Goodness Duru, Ghollah Kioko, Murhabazi Espoir, Elan van Biljon, Daniel Whitenack, Christopher
 1169 Onyefuluchi, Chris Chinenyere Emezie, Bonaventure F. P. Dossou, Blessing Sibanda, Blessing Bassey, Ayodele Olabiyi, Arshath Ramkilowan, Alp
 1170 Öktem, Adewale Akinfaderin, and Abdallah Bashir. 2020. Participatory Research for Low-resourced Machine Translation: A Case Study in African
 1171 Languages. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, Trevor Cohn, Yulan He, and Yang Liu (Eds.). Association for
 1172 Computational Linguistics, Online, 2144–2160. doi:10.18653/v1/2020.findings-emnlp.195
- 1173 [110] Toussaint Nothias. 2020. Access granted: Facebook's free basics in Africa. *Media, Culture & Society* 42, 3 (April 2020), 329–348. doi:10.1177/
 1174 0163443719890530
- 1175 [111] Rodney T. Ogawa and Betty Malen. 1991. Towards Rigor in Reviews of Multivocal Literatures: Applying the Exploratory Case Study Method.
 1176 *Review of Educational Research* 61, 3 (1991), 265–286. doi:10.3102/00346543061003265
- 1177 [112] Chinasa T. Okolo, Kehinde Aruleba, and George Obaido. 2023. *Responsible AI in Africa—Challenges and Opportunities*. Springer International
 1178 Publishing, Cham, 35–64. doi:10.1007/978-3-031-08215-3_3
- 1179 [113] Arnold Overwijk, Chenyan Xiong, Xiao Liu, Cameron VandenBerg, and Jamie Callan. 2022. ClueWeb22: 10 Billion Web Documents with Visual
 1180 and Semantic Information. (2022). doi:10.48550/ARXIV.2211.15848
- 1181 [114] Ciaran O'Faircheallaigh. 2015. Social Equity and Large Mining Projects: Voluntary Industry Initiatives, Public Regulation and Community
 1182 Development Agreements. *Journal of Business Ethics* 132, 1 (Nov. 2015), 91–103. doi:10.1007/s10551-014-2308-3
- 1183 [115] Arsenio Paez. 2017. Gray literature: An important resource in systematic reviews. *Journal of Evidence-Based Medicine* 10, 3 (2017), 233–240.
 1184 doi:10.1111/jebm.12266
- 1185 [116] Matthew J. Page, Joanne E. McKenzie, Patrick M. Bossuyt, Isabelle Boutron, Tammy C. Hoffmann, Cynthia D. Mulrow, Larissa Shamseer, Jennifer M.
 1186 Tetzlaff, Elie A. Akl, Sue E. Brennan, Roger Chou, Julie Glanville, Jeremy M. Grimshaw, Asbjorn Hróbjartsson, Manoj M. Lalu, Tianjing Li,
 1187 Elizabeth W. Loder, Evan Mayo-Wilson, Steve McDonald, Luke A. McGuinness, Lesley A. Stewart, James Thomas, Andrea C. Tricco, Vivian A.
 1188 Welch, Penny Whiting, and David Moher. 2021. The PRISMA 2020 statement: an updated guideline for reporting systematic reviews. *BMJ* 372
 1189 (March 2021), n71. doi:10.1136/bmj.n71
- 1190 [117] Joon Sung Park, Danielle Bragg, Ece Kamar, and Meredith Ringel Morris. 2021. Designing an Online Infrastructure for Collecting AI Data From
 1191 People With Disabilities. In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*. ACM, Virtual Event Canada,
 1192 52–63. doi:10.1145/3442188.3445870
- 1193 [118] Yong Jin Park and S. Mo Jones-Jang. 2023. Surveillance, Security, and Ai as Technological Acceptance. *AI and Society* 38, 6 (2023), 2667–2678.
 doi:10.1007/s00146-021-01331-9
- 1194 [119] Frank Pasquale and Haochen Sun. 2024. Consent and Compensation: Resolving Generative AI's Copyright Crisis. 4826695 (May 2024). doi:10.
 1195 2139/ssrn.4826695
- 1196

- [1197] [120] Michael Quinn Patton. 1991. Towards Utility in Reviews of Multivocal Literatures. *Review of Educational Research* 61, 3 (sept 1991), 287–292. doi:10.3102/00346543061003287
- [1198] [121] Charlotte Paul and Barbara Brookes. 2015. The Rationalization of Unethical Research: Revisionist Accounts of the Tuskegee Syphilis Study and the New Zealand “Unfortunate Experiment”. *American Journal of Public Health* 105, 10 (Oct. 2015), e12–19. doi:10.2105/AJPH.2015.302720
- [1199] [122] Guilherme Penedo, Quentin Malartic, Daniel Hesslow, Ruxandra Cojocaru, Alessandro Cappelli, Hamza Alobeidli, Baptiste Pannier, Ebtesam Almazrouei, and Julien Launay. 2023. The RefinedWeb Dataset for Falcon LLM: Outperforming Curated Corpora with Web Data, and Web Data Only. arXiv:2306.01116 (2023). doi:10.48550/arXiv.2306.01116 arXiv:2306.01116 [cs].
- [1200] [123] Maneesha Perera, Rajith Vidanaarachchi, Sangeetha Chandrashekeran, Melissa Kennedy, Brendan Kennedy, and Saman Halgamuge. 2025. Indigenous peoples and artificial intelligence: A systematic review and future directions. *Big Data & Society* 12, 2 (2025), 20539517251349170. doi:10.1177/20539517251349170
- [1201] [124] Claudio Santos Pinhanez and Edem Wornyo. 2025. Ethical Co-Development of AI Applications with Indigenous Communities. In *Proceedings of the Extended Abstracts of the CHI Conference on Human Factors in Computing Systems (CHI EA '25)*. Association for Computing Machinery, New York, NY, USA, 1–4. doi:10.1145/3706599.3706649
- [1202] [125] Ian Pool. 2016. *Colonialism's and postcolonialism's fellow traveller: the collection, use and misuse of data on indigenous people* (1st ed.). ANU Press. doi:10.22459/CAEPR38.11.2016.04
- [1203] [126] Rida Qadri, Michael Madaio, and Mary L. Gray. 2025. Confusing the Map for the Territory – Communications of the ACM. <https://cacm.acm.org/opinion/confusing-the-map-for-the-territory/>
- [1204] [127] Rida Qadri, Piotr Mirowski, and Remi Denton. 2025. AI and Non-Western Art Worlds: Reimagining Critical AI Futures through Artistic Inquiry and Situated Dialogue. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*. ACM, Yokohama Japan, 1–17. doi:10.1145/3706598.3714049
- [1205] [128] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. 2020. Exploring the limits of transfer learning with a unified text-to-text transformer. *J. Mach. Learn. Res.* 21, 1 (Jan. 2020), 140:5485–140:5551.
- [1206] [129] Jenalea Rajab, Anuoluwapo Aremu, Everlyn Asiko Chimoto, Dale Dunbar, Graham Morrissey, Fadel Thior, Luandrie Potgieter, Jessico Ojo, Atnafu Lambebo Tonja, Maushami Chetty, Wilhelmina NdapewaOnyothi Nekoto, Pelonomi Moiloa, Jade Abbott, Vukosi Marivate, and Benjamin Rosman. 2025. The Esethu Framework: Reimagining Sustainable Dataset Governance and Curation for Low-Resource Languages. arXiv:2502.15916 (2025). doi:10.48550/arXiv.2502.15916 arXiv:2502.15916 [cs].
- [1207] [130] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J Liu. [n. d.]. Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer. 2020 ([n. d.]).
- [1208] [131] Mohammad Rashidujjaman Rifat, Abdullah Hasan Safir, Sourav Saha, Jahedul Alam Junaed, Maryam Saleki, Mohammad Ruhul Amin, and Syed Ishtiaque Ahmed. 2024. Data, Annotation, and Meaning-Making: The Politics of Categorization in Annotating a Dataset of Faith-based Communal Violence. In *The 2024 ACM Conference on Fairness, Accountability, and Transparency*. ACM, Rio de Janeiro Brazil, 2148–2156. doi:10.1145/3630106.3659030
- [1209] [132] Katja Rogers, Teresa Hirzle, Sukran Karaosmanoglu, Paula Toledo Palomino, Ekaterina Durmanova, Seiji Isotani, and Lennart E. Nacke. 2024. An Umbrella Review of Reporting Quality in CHI Systematic Reviews: Guiding Questions and Best Practices for HCI. *ACM Trans. Comput.-Hum. Interact.* 31, 5 (Nov. 2024), 57:1–57:55. doi:10.1145/3685266
- [1210] [133] Caroline Running Wolf and Noelani Arista. 2020. *Indigenous Protocols in Action*. Indigenous Protocol and Artificial Intelligence Working Group and the Canadian Institute for Advanced Research, Honolulu, HI, 93–101. doi:10.11573/spectrum.library.concordia.ca.00986506
- [1211] [134] Nithya Sambasivan, Shivani Kapania, Hannah Highfill, Diana Akrong, Praveen Paritosh, and Lora M Aroyo. 2021. “Everyone wants to do the model work, not the data work”: Data Cascades in High-Stakes AI. In *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*. ACM, Yokohama Japan, 1–15. doi:10.1145/3411764.3445518
- [1212] [135] Advait Sarkar. 2023. Enough With “Human-AI Collaboration”. In *Extended Abstracts of the 2023 CHI Conference on Human Factors in Computing Systems*. ACM, Hamburg Germany, 1–8. doi:10.1145/3544549.3582735
- [1213] [136] Devansh Saxena, Ji-Youn Jung, Jodi Forlizzi, Kenneth Holstein, and John Zimmerman. 2025. AI Mismatches: Identifying Potential Algorithmic Harms Before AI Development. arXiv:2502.18682 (April 2025). doi:10.48550/arXiv.2502.18682 arXiv:2502.18682 [cs].
- [1214] [137] Morgan Klaus Scheuerman, Alex Hanna, and Remi Denton. 2021. Do Datasets Have Politics? Disciplinary Values in Computer Vision Dataset Development. *Proceedings of the ACM on Human-Computer Interaction* 5, CSCW2 (Oct. 2021), 1–37. doi:10.1145/3476058
- [1215] [138] Morgan Klaus Scheuerman, Kandrea Wade, Caitlin Lustig, and Jed R. Brubaker. 2020. How We've Taught Algorithms to See Identity: Constructing Race and Gender in Image Databases for Facial Analysis. *Proc. ACM Hum.-Comput. Interact.* 4, CSCW1 (May 2020), 58:1–58:35. doi:10.1145/3392866
- [1216] [139] Morgan Klaus Scheuerman, Katy Weathington, Tarun Mugunthan, Emily Denton, and Casey Fiesler. 2023. From Human to Data to Dataset: Mapping the Traceability of Human Subjects in Computer Vision Datasets. *Proceedings of the ACM on Human-Computer Interaction* 7, CSCW1 (April 2023), 1–33. doi:10.1145/3579488
- [1217] [140] Morgan Klaus Scheuerman, Allison Woodruff, and Jed R. Brubaker. 2025. How Data Workers Shape Datasets: The Role of Positionality in Data Collection and Annotation for Computer Vision. *Proc. ACM Hum.-Comput. Interact.* 9, 7 (Oct. 2025), CSCW300:1–CSCW300:42. doi:10.1145/3757481
- [1218] [141] Daniel Schiff, Bogdana Rakova, Aladdin Ayesh, Anat Fanti, and Michael Lennon. 2020. *Principles to Practices for Responsible AI: Closing the Gap*. doi:10.48550/arXiv.2006.04707 arXiv:2006.04707 [cs].
- [1219] [142] Manuscript submitted to ACM

- 1249 [142] Christoph Schuhmann, Romain Beaumont, Richard Vencu, Cade Gordon, Ross Wightman, Mehdi Cherti, Theo Coombes, Aarush Katta, Clayton
 1250 Mullis, Mitchell Wortsman, Patrick Schramowski, Srivatsa Kundurthy, Katherine Crowson, Ludwig Schmidt, Robert Kaczmarczyk, and Jenia Jitsev.
 1251 2022. LAION-5B: An open large-scale dataset for training next generation image-text models. (2022). doi:10.48550/ARXIV.2210.08402
- 1252 [143] Christoph Schuhmann, Richard Vencu, Romain Beaumont, Robert Kaczmarczyk, Clayton Mullis, Aarush Katta, Theo Coombes, Jenia Jitsev, and
 1253 Aran Komatsuaki. 2021. LAION-400M: Open Dataset of CLIP-Filtered 400 Million Image-Text Pairs. (2021). doi:10.48550/ARXIV.2111.02114
- 1254 [144] Eve Kosofsky Sedgwick. 1990. *Epistemology of the Closet*. University of California Press. Google-Books-ID: u5jgaOhhmpgC.
- 1255 [145] Judith Shapiro and John-Andrew McNeish. 2021. *Our Extractive Age: Expressions of Violence and Resistance* (1 ed.). Routledge, London. doi:10.4324/
 9781003127611
- 1256 [146] Renee Shelby, Shalaleh Rismani, Kathryn Henne, AJung Moon, Negar Rostamzadeh, Paul Nicholas, N'Mah Yilla, Jess Gallegos, Andrew Smart,
 1257 Emilio Garcia, and Gurleen Virk. 2023. Sociotechnical Harms of Algorithmic Systems: Scoping a Taxonomy for Harm Reduction. arXiv:2210.05791
 1258 (2023). <http://arxiv.org/abs/2210.05791> arXiv:2210.05791 [cs].
- 1259 [147] Jan Simson, Alessandro Fabris, and Christoph Kern. 2024. Lazy Data Practices Harm Fairness Research. In *The 2024 ACM Conference on Fairness,
 1260 Accountability, and Transparency*. ACM, Rio de Janeiro Brazil, 642–659. doi:10.1145/3630106.3658931
- 1261 [148] Shivalika Singh, Freddie Vargus, Daniel Dsouza, Börje F Karlsson, Abinaya Mahendiran, Wei-Yin Ko, Herumb Shandilya, Jay Patel, Deividas
 1262 Mataciunas, Laura OMahony, et al. 2024. Aya dataset: An open-access collection for multilingual instruction tuning. *arXiv preprint arXiv:2402.06619*
 1263 (2024).
- 1264 [149] Linda Tuhiwai Smith. 2021. *Decolonizing Methodologies: Research and Indigenous Peoples* (third edition ed.). Bloomsbury Publishing. Google-
 1265 Books-ID: EwA1EAAAQBAJ.
- 1266 [150] Gayatri Chakravorty Spivak. 1994. *Can the Subaltern Speak?* Routledge, London, 66–111.
- 1267 [151] Yolande Strengers, Jathan Sadowski, Zhuying Li, Anna Shimshak, and Florian “Floyd” Mueller. 2021. What Can HCI Learn from Sexual Consent?:
 1268 A Feminist Process of Embodied Consent for Interactions with Emerging Technologies. In *Proceedings of the 2021 CHI Conference on Human Factors
 in Computing Systems*. ACM, Yokohama Japan, 1–13. doi:10.1145/3411764.3445107
- 1269 [152] Harini Suresh and John V. Guttag. 2021. A Framework for Understanding Sources of Harm throughout the Machine Learning Life Cycle. In *Equity
 1270 and Access in Algorithms, Mechanisms, and Optimization*. 1–9. doi:10.1145/3465416.3483305 arXiv:1901.10002 [cs, stat].
- 1271 [153] Harini Suresh, Emily Tseng, Meg Young, Mary Gray, Emma Pierson, and Karen Levy. 2024. Participation in the age of foundation models. In
 1272 *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*. 1609–1621.
- 1273 [154] Alex S. Taylor. 2011. Out there. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI ’11)*. Association for
 1274 Computing Machinery, New York, NY, USA, 685–694. doi:10.1145/1978942.1979042
- 1275 [155] Jordan Taylor, Wesley Hanwen Deng, Kenneth Holstein, Sarah Fox, and Haiyi Zhu. 2024. Carefully Unmaking the “Marginalized User”: A Diffractive
 1276 Analysis of a Gay Online Community. *ACM Transactions on Computer-Human Interaction* (2024), 3673229. doi:10.1145/3673229
- 1277 [156] John Taylor and Tahu Kukutai (Eds.). 2016. *Indigenous Data Sovereignty: Toward an Agenda*. ANU Press, Acton, ACT, Australia. doi:10.22459/
 1278 CAEPR38.11.2016
- 1279 [157] Te Hiku Media. [n.d.]. *Kaitiakitanga License*. <https://github.com/TeHikuMedia/Kaitiakitanga-License> GitHub repository.
- 1280 [158] Jaime A. Teixeira da Silva. 2022. Handling Ethics Dumping and Neo-Colonial Research: From the Laboratory to the Academic Literature. *Journal
 1281 of Bioethical Inquiry* 19, 3 (2022), 433–443. doi:10.1007/s11673-022-10191-x
- 1282 [159] Tesfa-Alem Tekle. 2020. Refugees in Ethiopia’s camps raise privacy and exclusion concerns over UNHCR’s new digital registration. <https://globalvoices.org/2020/03/19/refugees-in-ethiopias-camps-raise-privacy-and-exclusion-concerns-over-unhcrs-new-digital-registration/>
- 1283 [160] Jim Thatcher, David O’Sullivan, and Dillon Mahmoudi. 2016. Data colonialism through accumulation by dispossession: New metaphors for daily
 1284 data. *Environment and Planning D: Society and Space* 34, 6 (Dec. 2016), 990–1006. doi:10.1177/0263775816633195
- 1285 [161] Scott Timcke. 2024. AI and the digital scramble for Africa. <https://roape.net/2024/07/11/ai-and-the-digital-scramble-for-africa/>
- 1286 [162] Attafu Lambebo Tonja, Bonaventure F. P. Dossou, Jessica Ojo, Jenalea Rajab, Fadel Thior, Eric Peter Wairagala, Anuoluwapo Aremu, Pelonomi
 1287 Moiloa, Jade Abbott, Vukosi Marivate, and Benjamin Rosman. 2024. InkubaLM: A small language model for low-resource African languages.
 arXiv:2408.17024 (2024). doi:10.48550/arXiv.2408.17024 arXiv:2408.17024 [cs].
- 1288 [163] Paola Tubaro, Antonio A. Casilli, Maxime Cornet, Clément Le Ludec, and Juana Torres Cierpe. 2025. Where does AI come from? A global case
 1289 study across Europe, Africa, and Latin America. (Feb. 2025). doi:10.1080/13563467.2025.2462137 arXiv:2502.04860 [cs].
- 1290 [164] Ding Wang, Shantanu Prabhakar, and Nithya Sambasivan. 2022. Whose AI Dream? In search of the aspiration in data annotation. arXiv:2203.10748
 1291 (March 2022). doi:10.48550/arXiv.2203.10748 arXiv:2203.10748 [cs].
- 1292 [165] Lining Wang, Vaishnav Kameswaran, and Hernisa Kacorri. 2025. Toward a Taxonomy of Algorithmic Harms for Disability: A Systematic Review.
 1293 *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society* 8, 3 (Oct. 2025), 2649–2665. doi:10.1609/aies.v8i3.36745
- 1294 [166] Karl Werder, Balasubramanian Ramesh, and Rongen (Sophia) Zhang. 2022. Establishing Data Provenance for Responsible Artificial Intelligence
 1295 Systems. *ACM Transactions on Management Information Systems* 13, 2 (2022), 1–23. doi:10.1145/3503488
- 1296 [167] Cedric Deslandes Whitney and Justin Norman. 2024. Real Risks of Fake Data: Synthetic Data, Diversity-Washing and Consent Circumvention. In
 1297 *The 2024 ACM Conference on Fairness, Accountability, and Transparency*. ACM, Rio de Janeiro Brazil, 1733–1744. doi:10.1145/3630106.3659002
- 1298 [168] David Gray Widder. 2024. Epistemic Power in AI Ethics Labor: Legitimizing Located Complaints. In *The 2024 ACM Conference on Fairness,
 1299 Accountability, and Transparency*. ACM, Rio de Janeiro Brazil, 1295–1304. doi:10.1145/3630106.3658973

- 1301 [169] David Gray Widder, Meredith Whittaker, and Sarah Myers West. 2024. Why ‘open’ AI systems are actually closed, and why this matters. *Nature*
1302 635, 8040 (Nov. 2024), 827–833. doi:10.1038/s41586-024-08141-1
- 1303 [170] Claes Wohlin. 2014. Guidelines for snowballing in systematic literature studies and a replication in software engineering. In *Proceedings*
1304 *of the 18th International Conference on Evaluation and Assessment in Software Engineering*. ACM, London England United Kingdom, 1–10.
1305 doi:10.1145/2601248.2601268
- 1306 [171] Meg Young, Upol Ehsan, Ranjit Singh, Emnet Tafesse, Michele Gilman, Christina Harrington, and Jacob Metcalf. 2024. Participation versus scale:
1307 Tensions in the practical demands on participatory AI. *First Monday* (April 2024). doi:10.5210/fm.v29i4.13642
- 1308 [172] Edibe Betul Yucer. 2025. AI is finally trying to speak African languages. Will this end a historic neglect? *TRT Global* (Aug. 2025). <https://trt.global/afrika-english/article/359e1362af39>
- 1309 [173] Dora Zhao, Morgan Klaus Scheuerman, Pooja Chitre, Jerome T. A. Andrews, Georgia Panagiotidou, Shawn Walker, Kathleen H. Pine, and Alice
1310 Xiang. 2024. A taxonomy of challenges to curating fair datasets. In *Proceedings of the 38th International Conference on Neural Information Processing*
1311 *Systems (NIPS ’24, Vol. 37)*. Curran Associates Inc., Red Hook, NY, USA, 97826–97858.
- 1312
1313
1314
1315
1316
1317
1318
1319
1320
1321
1322
1323
1324
1325
1326
1327
1328
1329
1330
1331
1332
1333
1334
1335
1336
1337
1338
1339
1340
1341
1342
1343
1344
1345
1346
1347
1348
1349
1350
1351
1352 Manuscript submitted to ACM

A Review Methodologies in HCI

Review methodologies in HCI face persistent challenges when systematic transparency must coexist with interpretive expertise and when evidence circulates across fragmented publication ecosystems [46, 132]. Systematic reviews establish protocols for organizing evidence within bounded domains, and PRISMA frameworks support reproducibility [116]. Both approaches reach limits in interdisciplinary settings where evidence types vary and epistemological frameworks conflict [6, 106].

Multivocal literature reviews (MLRs) offer one way to meet these demands and balance rigor with practical utility, as recent work in responsible AI demonstrates [92]. The approach originated in educational research as a methodological framework to impose systematic rigor on reviews of diverse documents [111], initiating a discussion that clarified that the standard of rigor must be situational and secondary to utility for practitioners [120]. Software engineering later adapted MLRs to capture both state-of-the-art research and state-of-practice knowledge [50, 51]. MLRs integrate peer-reviewed academic literature with grey literature such as organizational reports, policy documents, community statements, technical documentation, practitioner outputs, and multimedia materials. The widely cited Luxembourg definition characterizes grey literature as material produced by government, academia, industry, or community groups that is not controlled by commercial publishers [53]. Diversity of source types and timeliness are core advantages, since emerging practices often circulate outside formal publication channels and appear earlier than peer-reviewed work [115].

Credibility varies across grey-literature types, and assessments often depend on provenance, expertise, and recognized authority [72]. Multivocal approaches are particularly important for scholarship involving Indigenous and underserved communities. Lewis et al. [86] show that multivocality preserves heterogeneous viewpoints. The authors combine essays, protocols, and artistic works rather than imposing a single scholarly mode, a stance aligning with decolonizing methodologies that emphasize community-generated knowledge and Indigenous epistemic authority [9, 149]. Additionally, influential contributions in AI ethics and data governance often appear in organizational reports [24, 67], investigative journalism [61], public initiatives [1], and widely cited preprints [14]. MLR methods accommodate a diversity of distributed knowledge production and support synthesis across venues not fully captured by academic indexing.

B Search Strategy & Corpus Composition

Database searches were conducted iteratively between August 2024 and January 2025, complementing network referrals and citation snowballing. The final structured ACM Digital Library search was executed on January 31, 2025, using the advanced search interface with abstract and full-text indexing via personal subscription. Table 3 reports the four primary ACM query sets and their outcomes.

Aggregate results from the January 2025 ACM searches are summarized in Table 5. Across 1,914 hits, 1,201 items were screened, yielding 153 that met criteria and 48 unique sources after full-text review and duplicate removal. Similar comprehensive search strategies were applied to IEEE Xplore, ScienceDirect, Taylor & Francis Online, Wiley Online Library, Google Scholar, and Springer Link, following the same phased approach for queries.

Source Discovery.**C Corpus Creation Details****Datasheet Fields.**

Table 3. ACM Digital Library search queries and results.

Query	Exact search string	Results screened	Potentially relevant
Q1	((“data collection” OR “data production” OR “data curation” OR “dataset development”) AND (“artificial intelligence” OR “machine learning” OR “AI”) AND (“marginalized” OR “underrepresented” OR “underserved” OR “community” OR “indigenous”))	51 abstracts + 300 full-texts	45 (7 abstracts, 38 full-texts)
Q2	((“extractive” OR “exploitative” OR “data colonialism”) AND (“data practices” OR “dataset construction”) AND (“communities” OR “workers” OR “labor”))	3 abstracts + 182 full-texts	13 (1 abstract, 12 full-texts)
Q3	((“crowdsourcing” OR “platform labor”) AND (“bias” OR “fairness” OR “ethics”) AND (“marginalized” OR “vulnerable populations” OR “community harm”))	491 full-texts (200 screened)	32
Q4	((“participatory design” OR “community-led” OR “co-design”) AND (“ai development” OR “dataset creation”) AND (“sovereignty” OR “community engagement” OR “ethical data”))	122 full-texts	12

Table 4. Targeted ACM venue-specific searches.

Venue	Exact search string	Venue filter	Results summary
CHI Conference Proceedings	((“data collection” OR “data production” OR “data curation” OR “dataset development”) AND (“artificial intelligence” OR “machine learning” OR “AI”) AND (“marginalized” OR “underrepresented” OR “underserved” OR “community” OR “indigenous”))	CHI Conference on Human Factors in Computing Systems (all years)	296 hits; screened: first 200; potentially relevant: 15
FAccT Proceedings	((“data collection” OR “data production” OR “data curation” OR “dataset development”) AND (“artificial intelligence” OR “machine learning” OR “AI”) AND (“marginalized” OR “underrepresented” OR “underserved” OR “community” OR “indigenous”))	ACM Conference on Fairness, Accountability, and Transparency	143 hits; screened: all; potentially relevant: 37

Table 7. Description of datasheet fields

Column	Content
Identifier	In-line APA citation (author surname and year) used as a unique ID for tracking within the corpus.
APA Citation	Full APA reference for the source.
Title	Title of the publication or output.

Manuscript submitted to ACM

1457 1458 1459 1460 1461 1462 1463 1464 1465 1466 1467 1468 1469 1470 1471 1472 1473 1474 1475 1476 1477 1478 1479 1480 1481 1482 1483 1484 1485 1486 1487 1488 1489 1490 1491 1492 1493 1494 1495 1496 1497 1498 1499 1500 1501 1502 1503 1504 1505 1506 1507 1508	Column	Content
	Analytic Domain	Controlled list, multiple possible:
		<ul style="list-style-type: none"> • Data Relations: how data is scoped, justified, and negotiated • Data Labor: how curation work is arranged and carried out • Data Representation: how categories are constructed and based on what presences and absences • Data Infrastructure: how technical systems mediate data movement • Data Governance: how authority over data shapes downstream use
	Orientation	Controlled list:
		<ul style="list-style-type: none"> • Extractive: undermines consent, compensation, or benefit • High-Agency Principles: normative frameworks promoting stewardship, sovereignty, accountability • High-Agency Practices: operationalized, community-led, participatory, or sovereignty-based initiatives
	General Theme	Controlled list, multiple possible:
		<ul style="list-style-type: none"> • Community impacts and relations • Critical theory • Data labor • Data practices • Ethics frameworks
	Pipeline Stage	Specific process within a pipeline stage (controlled list):
		<ul style="list-style-type: none"> • Problem Understanding and Formulation • ML System Design and Development • Deployment and Impact • Cross-pipeline

Column	Content
Pipeline Sub-stage	<p>Specific process within a pipeline stage (controlled list):</p> <ul style="list-style-type: none"> • Institutional Prioritization and Funding • Product Conception and Design • Data Selection, Collection and Annotation • Model Architecture Selection and Design • Model Training and Evaluation • Product Testing • Product Launch • Cross-pipeline
Historical Era	<p>Era of data production practice (controlled list):</p> <ul style="list-style-type: none"> • Era 1: Curated datasets (pre-2009); no sources in corpus • Era 2: Crowdsourced benchmarks (2009–2017) • Era 3: Web-scraped/foundation models (2017–present) • Multi-era: Spans multiple eras or provides historical analysis
Primary Pattern(s) / Pathway(s)	<p>The specific extractive or high-agency behavior described in the source. Between one and three tags were assigned per source in order of relevance. For sources that provide conceptual, historical, or framing contributions without mapping directly onto an identified pattern, we assigned Other/NA (conceptual framing).</p>
Triangle Coverage	<p>Engagement with the three scoping domains that defined corpus eligibility:</p> <ul style="list-style-type: none"> • A – AI contexts • D – Data production practices • C – Community impacts <p>Because community impacts (C) establish the outer bounds of the review, included sources substantively address all three domains, though with varying emphases. Codes (A, D, C or combinations ADC, DC, AD) indicate which domains are explicitly developed within the source. An accompanying Rationale column explains the basis for inclusion and the specific ways each source engages A/D/C beyond passing mention.</p>
How Source Was Found	<p>White literature (journal papers, conference proceedings, books) or Grey literature (reports, policy documents, theses, community outputs, blogs).</p>
Keywords	<p>3–5 terms for coding/search, ordered Geography → Data/technical → Community/impact.</p>

Column	Content
Geographic Region of Focus	Region or community under study (controlled list): Africa, APAC, EU/UK, LatAm, MENA, North America, Oceania, Multiple regions, Not regionally specific (globally framed advocacy, transnational collectives, or technical works not tied to one region).
Author Affiliations	High-level institutional grouping of authors. If multiple affiliations, code majority grouping here; record full details in Authorship & Positionality Context. Controlled list: Academic; Government; Industry; NGO/Non-profit; Mixed; Journalist/Other/Not sure
Geographic Area of Author(s)	Full institution name and country of the lead author(s)
Institution	Region of lead author's institution (controlled list, same regions as above).
Authorship and Positionality Context	Complete authorship profile, including all institutions, geographic distribution, equal contribution notes, and any relevant statements on positionality or disciplinary traditions.
Summary	≤ 120-word synopsis. Structure: Topic → Method → Findings → Link to AI data production + community impacts.

Table 5. ACM search results summary.

Category	Count
Total primary searches (query sets)	6
Total venue-specific searches	2
Total hits across all searches	1,914
Total items screened (varied by search size)	1,201
Items meeting inclusion criteria after screening	153
Items retained after full-text review	89
Final unique sources for corpus (after duplicate removal)	48

Table 6. Discovery method distribution (N=350 sources).

Method	Sources	Percent
Database searches	174	50%
Existing networks/organizations	73	21%
Citation snowballing	51	15%
Iterative keyword search	31	9%
Hand-searching journals	21	6%
Total	350	100%

Corpus Summary.

Received 4 September 2025; revised 4 December 2025; accepted 5 June 2009

Table 8. Corpus composition summary (N=350 sources)

Category	Sub-category	Count (%)
Orientation	Extractive Practices	141 (41%)
	High-Agency Principles	116 (33%)
	High-Agency Practices	93 (27%)
Source Type	White literature	258 (74%)
	Grey literature	92 (26%)
Geographic Focus	Not regionally specific	150 (43%)
	Multiple areas	61 (17%)
	North America	41 (11%)
	Africa	38 (11%)
	Oceania	19 (5%)
	APAC	15 (4%)
	EU/UK	11 (3%)
	LatAm	12 (3%)
	MENA	3 (1%)
Author Affiliation (lead only)	Academic	182 (52%)
	Mixed	96 (27%)
	Industry	37 (11%)
	NGO/Non-profit	20 (6%)
	Journalist/Other	11 (3%)
Pipeline Stage	Government	4 (1%)
	ML System Design & Development	183 (52%)
	Problem Understanding & Formulation	90 (25%)
	Cross-pipeline	55 (16%)
Historical Era	Deployment & Impact	22 (6%)
	Era 3 (2017–present)	241 (69%)
	Multi-era	95 (27%)
	Era 2 (2009–2017)	14 (4%)
	Era 1 (pre-2009)	0 (0%)